

XcalableMP Programming Model and Language

Hitoshi Murai, Masahiro Nakao, and Mitsuhsa Sato

Abstract XcalableMP (XMP) is a directive-based language extension of Fortran and C for distributed-memory parallel computers, and can be classified as a partitioned global address space (PGAS) language. One of remarkable characteristics of XMP is that it supports both of global-view and local-view parallel programming. This chapter describes the programming model and language specification of XMP.

1 Introduction

Distributed-memory systems are generally used for large-scale simulations. To program such systems, Message Passing Interface (MPI) is widely adopted. However, programming with MPI is difficult because programmers must describe inter-process communications with consideration of the execution flow of their programs, which might cause deadlocks or wrong results.

To address this issue, a parallel language named High Performance Fortran (HPF) was proposed in 1991. With HPF, programmers can execute their serial programs in parallel by inserting minimal directives into them. If the programmers specify data distribution with HPF directives, the compilers do all other tasks for parallelization (e.g. communication generation and work distribution). However, HPF was not widely accepted eventually because the compilers' automatic processing prevents

Hitoshi Murai
RIKEN Center for Computational Science, 7-1-26 Minatojima-minami-machi, Chuo-ku, Kobe,
Hyogo 650-0047, Japan, e-mail: h-murai@riken.jp

Masahiro Nakao
RIKEN Center for Computational Science, 7-1-26 Minatojima-minami-machi, Chuo-ku, Kobe,
Hyogo 650-0047, Japan, e-mail: masahiro.nakao@riken.jp

Mitsuhsa Sato
RIKEN Center for Computational Science, 7-1-26 Minatojima-minami-machi, Chuo-ku, Kobe,
Hyogo 650-0047, Japan, e-mail: msato@riken.jp

the programmers from performance tuning, and the performance depends heavily on the environment (e.g. compiler and hardware)

Note: For more detail, please refer: Ken Kennedy, Charles Koelbel and Hans Zima: The Rise and Fall of High Performance Fortran: An Historical Object Lesson, Proc. 3rd ACM SIGPLAN History of Programming Languages Conf. (HOPL-III), pp. 7-17-22 (2007).

In such circumstance, to develop a new parallel programming model that enables easy parallelization of existing serial programs and design a new language based on it, “the XMP Specification Working Group” was established in 2008. This group utilized the lessons from the experience of HPF to define a new parallel language *XcalableMP* (XMP). The group was reorganized to one of the working groups of PC Cluster Consortium in 2011.

It is learned from the lessons of HPF that more automatic processing of compilers increases the gap between a program and its execution, and, as a result, decreases the usability of the language. In XMP, the programmers specify explicitly the details of parallel programs on the basis of compiler directives to make their execution easy-to-understand. In particular, they can specify explicitly communication, synchronization, data mapping, and work mapping to facilitate performance tuning. In addition, XMP supports features for one-sided communication on each process, which was not available in HPF. This feature might enable programmers to implement parallel algorithms more easily.

In this chapter, an overview of the programming model and language specification of XMP is shown. You can find the latest and complete language specification of XMP in: XcalableMP Specification Working Group, XcalableMP Specification Version 1.4, <http://xcalablemp.org/download/spec/xmp-spec-1.4.pdf> (2018).

1.1 Target Hardware

The target of XcalableMP is distributed-memory multicomputers (Fig. 1). Each compute node, which may contain several cores, has its own local memory (shared by the cores, if any), and is connected with the others via an interconnection network. Each node can access its local memory directly and remote memory (the memory of another node) indirectly (i.e. via inter-node communication). However, it is assumed that accessing remote memory may be much slower than accessing local memory.




Fig. 1 Target hardware of XMP.

1.2 Execution Model

An XcalableMP program execution is based on the Single Program Multiple Data (SPMD) model, where each node starts execution from the same main routine, and continues to execute the same code independently (i.e. asynchronously) until it encounters an XcalableMP construct (Fig. 2).

A set of nodes that executes a procedure, statement, loop, a block, etc. is referred to as its *executing node set*, and is determined by the innermost `task`, `loop`, or `array` directive surrounding it dynamically, or at runtime. The *current executing node set* is an executing node set of the current context, which is managed by the XcalableMP runtime system on each node.

The current executing node set at the beginning of the program execution, or *entire node set*, is a node set that contains all the available nodes, which can be specified in an implementation-defined way (e.g. through a command-line option).

When a node encounters at runtime either a `loop`, `array`, or `task` construct, and is contained by the node set specified (explicitly or implicitly) by the `on` clause of the directive, it updates the current executing node set with the specified one and executes the body of the construct, after which it resumes the last executing node set and proceeds to execute the subsequent statements.

In particular, when a node in the current executing node set encounters a `loop` or an `array` construct, it executes the loop or the array assignment in parallel with the other nodes, so that each iteration of the loop or element of the assignment is independently executed by the node in which the specified data element resides.

When a node encounters a synchronization or a communication directive, synchronization or communication occurs between it and the other nodes. That is, such

global constructs are performed collectively by the current executing nodes. Note that neither synchronization nor communication occurs unless these constructs are specified.




Fig. 2 Execution model of XMP.

1.3 Data Model

There are two classes of data in XcalableMP: *global data* and *local data*. Data declared in an XcalableMP program are local by default.

Global data are distributed onto a node set by the `align` directive (see section ??). Each fragment of distributed global data is allocated in the local memory of a node in the node set.

Local data comprises all data that are not global. They are replicated within the local memory of each of the executing nodes.

A node can access directly only local data and sections of global data that reside in its local memory. To access data in remote memory, explicit communication must be specified in such ways as global communication constructs and coarray assignments.

1.4 Programming Models

1.4.1 Partitioned Global Address Space

XMP can be classified as a *partitioned global address space (PGAS)* language, such as Co-Array Fortran [1], Unified Parallel C [2], and Chapel [3].

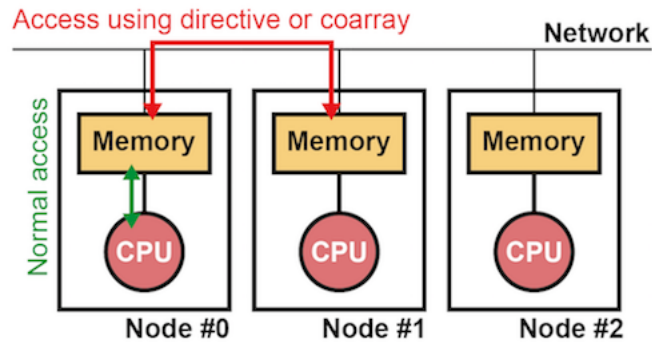


Fig. 3 Data model of XMP.

In such PGAS languages, multiple executing entities (i.e. threads, processes, or nodes in XMP) share a part of their address space, which is, however, partitioned and a portion of which is local to each executing entity.

The two programming models, global-view and local-view, that XMP supports to achieve high performance and productivity on PGAS are explained below.

1.4.2 Global-view Programming Model

The global-view programming model is useful when, starting from a serial version of a program, the programmer parallelizes it in a data-parallel style by adding directives with minimum modification. Based on this model, the programmer specifies the distribution of data among nodes using the data distribution directives. The loop construct assigns each iteration of a loop to the node at which the computed data is located. The global-view communication directives are used to synchronize nodes, maintain the consistency of shadow areas of distributed data, and move sections of distributed data globally. Note that the programmer must specify explicitly communication to make all data references in their program local using appropriate directives.

In many cases, the XcalableMP program following to the global-view programming model is based on a serial program, and it can produce the same result, regardless of the number of nodes (Fig. 4).

There are three groups of directives for this model:

- *Data mapping*, which specifies the data distribution and mapping to nodes
- *Work mapping (parallelization)*, which specifies the work distribution and mapping to nodes.
- *Communication and synchronization*, which specify how a node communicates and synchronizes with the other nodes.

Because these directives are ignored as a comment by the compilers of base languages (Fortran and C), an XcalableMP program can usually be compiled by them to ensure that they run properly.

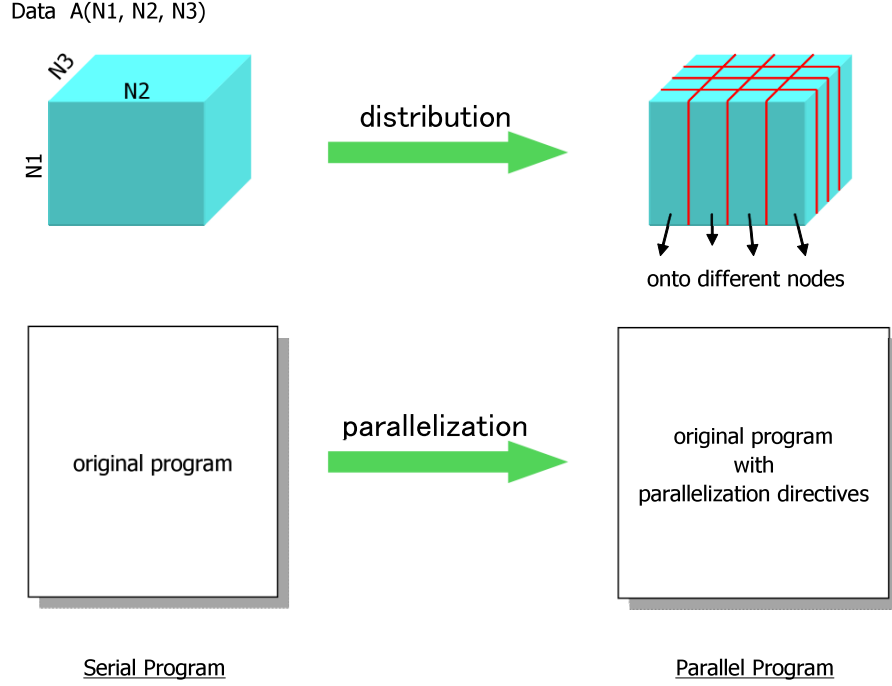


Fig. 4 Parallelization based on the global-view programming model.

1.4.3 Local-view Programming Model

The local-view programming model is suitable for programs that implement an algorithm and a remote data reference that are to be executed by each node (Fig. 5).

For this model, some language extensions and directives are provided. The coarray notation, which is imported from Fortran 2008, is one such extension, and can be used to explicitly specify data on which node is to be accessed. For example, the expression of $A(i)[N]$ is used to access an array element of $A(i)$ located on the node N . If the access is a reference, then a one-sided communication to read the value from the remote memory (i.e. the *get* operation) is issued by the executing node. If the access is a definition, then a one-sided communication to write the value to the remote memory (i.e. the *put* operation) is issued by the executing node.

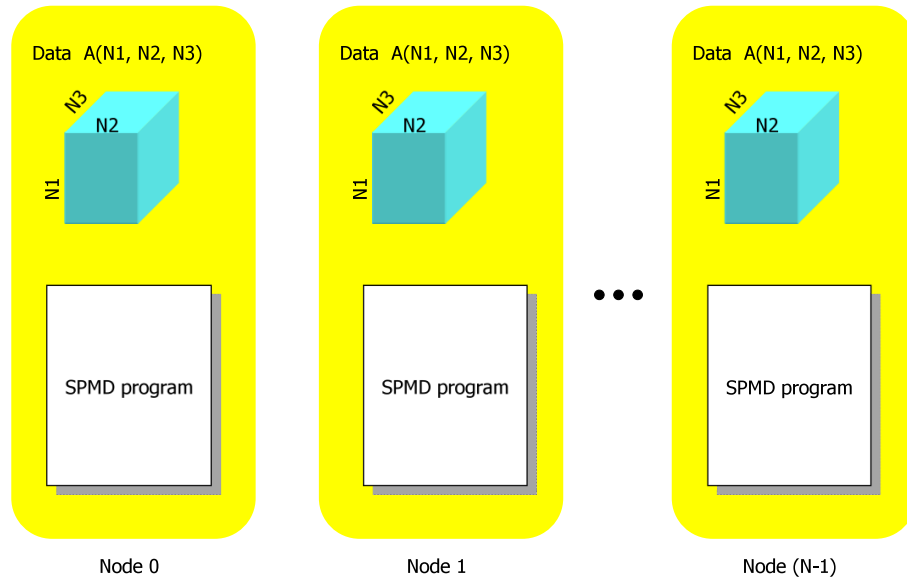


Fig. 5 Local-view programming model.

1.4.4 Mixture of Global View and Local View

In the global-view model, nodes are used to distribute data and works. In the local-view model, nodes are used to address remote data in the coarray notation. In application programs, the programmers should choose an appropriate data model according to the characteristics of their program. Fig. 6 illustrates the global view and the local view of data.

Data can have both a global view and a local view, and can be accessed in both of the views. XcalableMP provides a directive to give the local name (alias) to global data declared in the global-view programming model to enable them to also be accessed in the local-view programming model. This feature is useful to optimize a certain part of a program by using explicit remote data access in the local-view programming model.

1.5 Base Languages

The XcalableMP language specification is defined on the basis of Fortran and C as the base languages. More specifically, the base language of XcalableMP Fortran is Fortran 90 or later, and that of XcalableMP C is ISO C90 (ANSI C89) or later with some extensions (see below).

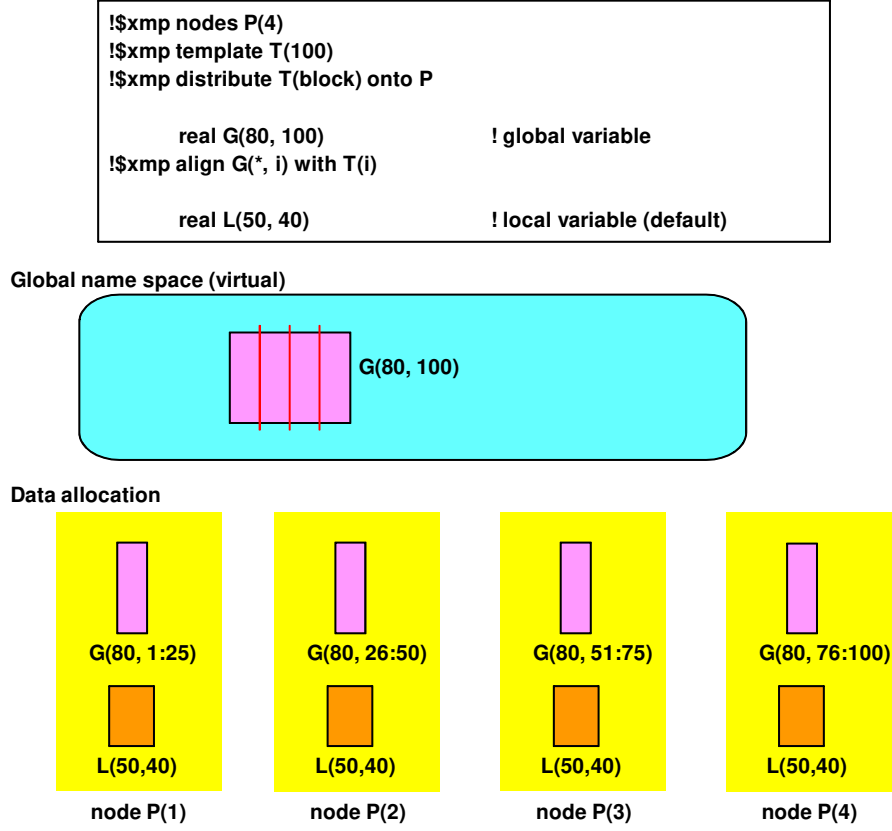


Fig. 6 Global view and local view.

1.5.1 Array Section in XcalableMP C

In XcalableMP C, the base language C is extended so that a part of an array, that is, an *array section* or *subarray*, can be put in an *array assignment statement*, which is described in 1.5.2, and some XcalableMP constructs. An array section is built from a subset of the elements of an array, which is specified by a sequence of square-bracketed integer expressions or *triplets*, which are in the form of:

$$[\textit{base}] : [\textit{length}] [: \textit{step}]$$

When *step* is positive, the *triplet* specifies a set of subscripts that is a regularly spaced integer sequence of length *length* beginning with *base* and proceeding in increments of *step* up to the largest. The same applies to negative *step* too.

When *base* is omitted, it is assumed to be 0. When *length* is omitted, it is assumed to be the number of remainder elements of the dimension of the array. When *step* is omitted, it is assumed to be 1.

Assuming that an array A is declared by the following statement,


```
int A[100];
```

some array sections can be specified as follows:

```
A[10:10]  array section of 10 elements from A[10] to A[19]
A[10:]    array section of 90 elements from A[10] to A[99]
A[:10]    array section of 10 elements from A[0] to A[9]
A[10:5:2] array section of 5 elements from A[10] to A[18] by step 2
A[:]      array section of the whole of A
```

1.5.2 Array Assignment Statement in XscalableMP C

In XscalableMP C, the base language C is also extended so that it supports array assignment statements just as Fortran does.

With such statement, the value of each element of the result of the right-hand side expression is assigned to the corresponding element of the array section on the left-hand side. When an operator or an elemental function is applied to array sections in the right-hand side expression, it is evaluated to an array section that has the same shape as that of the operands or arguments, and each element of which is the result of the operator or function applied to the corresponding element of the operands or arguments. A scalar object is assumed to be an array section that has the same shape as that of the other array section(s) in the expression or on the left-hand side, and where each element has its value.

Note that an array assignment is a statement, and therefore cannot appear as an expression in any other statements.

An array assignment statement in the fourth line copies the five elements from B[0] to B[4] into the elements from A[5] to A[9].

XscalableMP C
<pre>int A[10]; int B[5]; ... A[5:5] = B[0:5];</pre>

1.6 Interoperability

Most of existing parallel applications are written with MPI. It is not realistic to port them over to XMP because each of them consists of millions of lines.

Because XMP is interoperable with MPI, users can develop an XMP application by modifying a part of an existing one instead of rewriting it totally. Besides, when developing a parallel application from scratch, it is possible to use XMP to write a complicated part of, for example, domain decomposition while they use MPI, which

could be faster than XMP, to write a hot-spot part that need to be tuned carefully. In addition, XMP is interoperable with OpenMP and Python (see Chapter ??).

It might be difficult to develop an application with just one programming language or framework since it generally has its own strong and weak points. Thus, an XMP program is interoperable with those in other languages to provide both high productivity and performance.

2 Data Mapping

2.1 nodes Directive

The `nodes` directive declares a *node array*, which is an array-like arrangement of nodes in a node set. A node array can be multi-dimensional.

	XcalableMP C	
<code>#pragma xmp nodes p[4]</code>		
<div style="text-align: center; border-bottom: 1px solid black;">XcalableMP Fortran</div> <p><code>!\$xmp nodes p(4)</code></p>		

The `nodes` directive declares a one-dimensional node array `p` that includes four nodes. In XMP/C, it is zero-based and consists of `p[0]`, `p[1]`, `p[2]`, and `p[3]`. In XMP/Fortran, it is one-based and consists of `p(1)`, `p(2)`, `p(3)`, and `p(4)`.

	XcalableMP C	
<code>#pragma xmp nodes p[2][3]</code>		
<div style="text-align: center; border-bottom: 1px solid black;">XcalableMP Fortran</div> <p><code>!\$xmp nodes p(3,2)</code></p>		

The `nodes` directive declares two-dimensional node array `p` that includes six nodes. In XMP/C, it consists of `p[0][0]`, `p[0][1]`, `p[0][2]`, `p[1][0]`, `p[1][1]`, and `p[1][2]`. In XMP/Fortran, it consists of `p(1,1)`, `p(2,1)`, `p(3,1)`, `p(1,2)`, `p(2,2)`, and `p(3,2)`.

Note: The ordering of the elements in a node array follows that of a normal array in the base language, C or Fortran.

	XcalableMP C	
<code>#pragma xmp nodes p[*]</code>		
<div style="text-align: center; border-bottom: 1px solid black;">XcalableMP Fortran</div> <p><code>!\$xmp nodes p(*)</code></p>		

An asterisk can be specified as the size in the `nodes` directive to declare a *dynamic* node array. In the above code, one-dimensional dynamic node array `p` is declared with an asterisk as the size. The actual size of a dynamic node array is determined at runtime to fit the size of the current executing node set. For example, when the programmer runs the sample code with three nodes, the node array `p` includes three nodes.

They can also declare multi-dimensional dynamic node arrays with an asterisk.

XcalableMP C
#pragma xmp nodes p[*][3]
XcalableMP Fortran
!\$xmp nodes p(3,*)

When the programmer runs the sample code with 12 nodes, the node array `p` has a shape of 4x3, in C, or 3x4, in Fortran.

Note: The user can put an asterisk only in the last dimension, in XMP/Fortran, or the first dimension, in XMP/C of the node array.

Hint: The dynamic node array may interfere with compiler optimizations. In general, programs with static ones achieve better performance.

The programmer can declare a node subarray derived from an existing node array. Node subarrays can be used, for example, to optimize inter-node communication by reducing the number of nodes participating in the communication.

XcalableMP C
#pragma xmp nodes p[16] #pragma xmp nodes q[8]=p[0:8] #pragma xmp nodes r[4][2]=p[8:8]
XcalableMP Fortran
!\$xmp nodes p(16) !\$xmp nodes q(8)=p(1:8) !\$xmp nodes r(2,4)=p(9:16)

In line 1, a node array `p` including 16 nodes is declared. In line 2, a node subarray `q` corresponding to the first half of `p` is declared. In line 3, a two-dimensional node subarray `r` corresponding to the latter half of `p` is declared.

The programmer can declare a *n*-dimensional node subarray derived from a *m*-dimensional one.

```

XcalableMP C
#pragma xmp nodes p[4][2]
#pragma xmp nodes row[4]=p[*][:]
#pragma xmp nodes col[2]=p[*][:]

XcalableMP Fortran
!$xmp nodes p(2,4)
!$xmp nodes row(4)=p(*,:)
!$xmp nodes col(2)=p(:,*)

```

In line 1, a two-dimensional node array p including 4×2 nodes is declared. In line 2, a node subarray row derived from a single row of p is declared. In line 3, a node subarray col derived from a single column of p is declared.

A colon represents a triplet which indicate all possible indices in the dimension. An asterisk indicate the index of the current executing node in the dimension. For example, $col[2]$ corresponds to $p[0][0:2]$ on nodes $p[0][0]$ and $p[0][1]$, and to $p[1][0:2]$ on nodes $p[1][0]$ and $p[1][1]$ in XMP/C. Similarly, $col(2)$ corresponds to $p(1:2, 1)$ on nodes $p(1, 1)$ and $p(2, 1)$, and to $p(1:2, 2)$ on nodes $p(1, 2)$ $p(2, 2)$ in XMP/Fortran.

XMP/C

$p[0][0]$	$p[0][1]$	$p[0][0]$	$p[0][1]$	$col[2]$
$p[1][0]$	$p[1][1]$	$p[1][0]$	$p[1][1]$	$col[2]$
$p[2][0]$	$p[2][1]$	$p[2][0]$	$p[2][1]$	$col[2]$
$p[3][0]$	$p[3][1]$	$p[3][0]$	$p[3][1]$	$col[2]$
$row[4]$	$row[4]$			

XMP/Fortran

$p(1,1)$	$p(2,1)$	$p(1,1)$	$p(2,1)$	$col[2]$
$p(1,2)$	$p(2,2)$	$p(1,2)$	$p(2,2)$	$col[2]$
$p(1,3)$	$p(2,3)$	$p(1,3)$	$p(2,3)$	$col[2]$
$p(1,4)$	$p(2,4)$	$p(1,4)$	$p(2,4)$	$col[2]$
$row[4]$	$row[4]$			

Fig. 7 Node subarrays.

In XMP/C, `row[0]` corresponds to `p[0][0]` and `p[0][1]` on `p[:,0]` and `p[:,1]`, respectively; `col[0]` corresponds to `p[0][0]`, `p[1][0]`, `p[2][0]`, and `p[3][0]` on `p[0,:]`, `p[1,:]`, `p[2,:]`, `p[3,:]`, respectively. In XMP/Fortran, `row(1)` corresponds to `p(1,1)` and `p(2,1)` on `p(1,:)` and `p(2,:)`, respectively; `col(1)` corresponds to `p(1,1)`, `p(1,2)`, `p(1,3)`, and `p(1,4)` on `p(:,1)`, `p(:,2)`, `p(:,3)`, `p(:,4)`, respectively.

Note: The semantics of an asterisk in a node reference is different from that in a declaration.

2.2 template Directive

The `template` directive declares a *template*, which is a virtual array that is used as a “template” of parallelization in the programs and to be distributed onto a node array.

XcalableMP C
<code>#pragma xmp template t[10]</code>

XcalableMP Fortran
<code>!\$xmp template t(10)</code>

This `template` directive declares a one-dimensional template `t` having ten elements. Templates are indexed in the similar manner to arrays in the base languages. For the above examples, the template `t` is indexed from zero to nine (i.e. `t[0] ... t[9]`), in XMP/C, or one to ten (i.e. `t(1) ... t(10)`), in XMP/Fortran.

Hint: In many cases, a template should be declared to have the same shape as your target array.

XcalableMP C
<code>#pragma xmp template t[10][20]</code>

XcalableMP Fortran
<code>!\$xmp template t(20,10)</code>

The `template` directive declares a two-dimensional template `t` that has 10x20 elements. In XMP/C, `t` is indexed from `t[0][0]` to `t[9][19]`, and, in XMP/Fortran, from `t(1,1)` to `t(20,10)`.

XcalableMP C
<code>#pragma xmp template t[:]</code>

XcalableMP Fortran

```
!$xmp template t(:)
```

In the above examples, a colon instead of an integer is specified as the size to declare a one-dimensional dynamic template `t`. The colon indicates that the size of the template is not fixed and to be fixed at runtime by the `template_fix` construct (Sec. 2.6).

2.3 distribute Directive

The `distribute` directive specifies a distribution of the target template. Either of *block*, *cyclic*, *block-cyclic*, or *gblock* (i.e. uneven block) can be specified to distribute a dimension of a template.

2.3.1 Block Distribution

XcalableMP C

```
#pragma xmp distribute t[block] onto p
```

XcalableMP Fortran

```
!$xmp distribute t(block) onto p
```

The target template `t` is divided into contiguous blocks and distributed among nodes in the node array `p`. Let's suppose that the size of the template is N and the number of nodes is K . If N is divisible by K , a block of size N/K are assigned to each node; otherwise, a block of size $\text{ceil}(N/K)$ is assigned to each of $N/\text{ceil}(N/K)$ nodes, a block of size $\text{mod}(N, K)$ to one node, and no block to $(K - N/\text{ceil}(N/K) - 1)$ nodes. The block distribution is useful for regular computations such as a stencil one.

Note: The function $\text{ceil}(x)$ returns a minimum integer value greater than x , and $\text{mod}(x, y)$ returns x modulo y .

XcalableMP C

```
#pragma xmp nodes p[3]
#pragma xmp template t[22]
#pragma xmp distribute t[block] onto p
```

XcalableMP Fortran

```
!$xmp nodes p(3)
!$xmp template t(22)
!$xmp distribute t(block) onto p
```




Fig. 8 Block distribution.

Since $\text{ceil}(22/3)$ is 8, eight elements are allocated on each of $p[0]$ and $p[1]$, and the remaining six elements are allocated on $p[2]$.

2.3.2 Cyclic Distribution

```

XcalableMP C
#pragma xmp distribute t[cyclic] onto p

XcalableMP Fortran
!$xmp distribute t(cyclic) onto p

```

The target template t is divided into chunks of size one and distributed among nodes in the node array p in a round-robin manner. The cyclic distribution is useful for the case where the load on each element of the template is not balanced.

```

XcalableMP C
#pragma xmp nodes p[3]
#pragma xmp template t[22]
#pragma xmp distribute t[cyclic] onto p

XcalableMP Fortran
!$xmp nodes p(3)
!$xmp template t(22)
!$xmp distribute t(cyclic) onto p

```



Fig. 9 Cyclic distribution.

2.3.3 Block-cyclic Distribution

XcalableMP C

```
#pragma xmp distribute t[cyclic(w)] onto p
```

XcalableMP Fortran

```
!$xmp distribute t(cyclic(w)) onto p
```

The target template t is divided into chunks of size w and distributed among nodes in the node array p in a round-robin manner. The block-cyclic distribution is useful for the case where the load on each element of the template is not balanced but the locality of the elements is required.


XcalableMP C

```
#pragma xmp nodes p[3]
#pragma xmp template t[22]
#pragma xmp distribute t[cyclic(3)] onto p
```

XcalableMP Fortran

```
!$xmp nodes p(3)
!$xmp template t(22)
!$xmp distribute t(cyclic(3)) onto p
```

XMP/C



XMP/Fortran





Fig. 10 Block-cyclic distribution.

2.3.4 Gblock Distribution

XcalableMP C

```
#pragma xmp distribute t[gblock(W)] onto p
```

XcalableMP Fortran

```
!$xmp distribute t(gblock(W)) onto p
```

The target template t is divided into contiguous blocks of size $W[0], W[1], \dots$, in XMP/C, or $W(1), W(2), \dots$, in XMP/Fortran, and distributed among nodes in the node array p . The array W is called a mapping array. The programmer can specify irregular (uneven) block distribution with the gblock format.


```

XcalableMP C
#pragma xmp nodes p[3]
#pragma xmp template t[22]
int W[3] = {6, 11, 5};
#pragma xmp distribute t[gblock(W)] onto p

XcalableMP Fortran
!$xmp nodes p(3)
!$xmp template t(22)
integer, parameter :: W(3) = (/6,11,5/)
!$xmp distribute t(gblock(W)) onto p

```




Fig. 11 Gblock distribution.

The programmer can specify an asterisk instead of a mapping array in the gblock distribution to defer fixing the actual distribution. In such a case, the actual distribution will be fixed at runtime by using `template_fix` construct.

2.3.5 Distribution of Multi-dimensional Templates

The programmer can distribute a multi-dimensional template onto a node array.

```

XcalableMP C
#pragma xmp nodes p[2][2]
#pragma xmp template t[10][10]
#pragma xmp distribute t[block][block] onto p

XcalableMP Fortran
!$xmp nodes p(2,2)
!$xmp template t(10,10)
!$xmp distribute t(block,block) onto p

```

The `distribute` directive declares the distribution of a two-dimensional template `t` onto a two-dimensional node array `p`. Each dimension of the template is divided in a block manner and each of the rectangular region is assigned to a node.

The programmer can specify a different distribution format in each of the dimension of a template.




Fig. 12 Example of multi-dimensional distribution (1).

```

XcalableMP C
#pragma xmp nodes p[2][2]
#pragma xmp template t[10][10]
#pragma xmp distribute t[block][cyclic] onto p

XcalableMP Fortran
!$xmp nodes p(2,2)
!$xmp template t(10,10)
!$xmp distribute t(cyclic,block) onto p

```

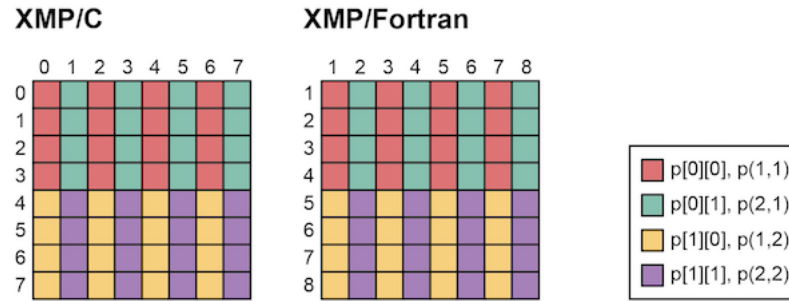


Fig. 13 Example of multi-dimensional distribution (2).

When an asterisk is specified in a `distribute` directive as a distribution format, the target dimension is “non-distributed.” In the following example, the first dimension is distributed in a block manner and the second dimension is non-distributed.

```

XcalableMP C
#pragma xmp nodes p[4]
#pragma xmp template t[10][10]
#pragma xmp distribute t[block][*] onto p

```

```

XcalableMP Fortran
!$xmp nodes p(4)
!$xmp template t(10,10)
!$xmp distribute t(*,block) onto p

```




Fig. 14 Example of multi-dimensional distribution (3).

2.4 align Directive

The `align` directive specifies that an array is to be mapped in the same way as a specified template. In other words, an `align` directive defines the correspondence of elements between an array and a template, and each of the array element is allocated on the node where the corresponding template element is allocated.

```

XcalableMP C
#pragma xmp nodes p[4]
#pragma xmp template t[8]
#pragma xmp distribute t[block] onto p
int a[8];
5 #pragma xmp align a[i] with t[i]

XcalableMP Fortran
!$xmp nodes p(4)
!$xmp template t(8)
!$xmp distribute t(block) onto p
integer :: a(8)
5 !$xmp align a(i) with t(i)

```

The array `a` is decomposed and laid out so that each element `a(i)` is colocated with the corresponding template element `t(i)`.

The `align` directive can also be used for multi-dimensional arrays.

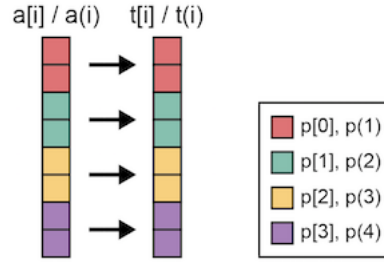


Fig. 15 Example of array alignment (1).

```

XcalableMP C
#pragma xmp nodes p[2][2]
#pragma xmp template t[8][8]
#pragma xmp distribute t[block][block] onto p
int a[8][8];
5 #pragma xmp align a[i][j] with t[i][j]

XcalableMP Fortran
!$xmp nodes p(2,2)
!$xmp template t(8,8)
!$xmp distribute t(block,block) onto p
integer :: a(8,8)
5 !$xmp align a(j,i) with t(j,i)

```




Fig. 16 Example of array alignment (2).

The programmer can align an n -dimensional array with an m -dimensional template for $n > m$.

```

XcalableMP C
#pragma xmp nodes p[4]
#pragma xmp template t[8]
#pragma xmp distribute t[block] onto p

```

```

int a[8][8];
5 #pragma xmp align a[i][*] with t[i]

```

XscalableMP Fortran

```

!$xmp nodes p(4)
!$xmp template t(8)
!$xmp distribute t(block) onto p
integer :: a(8,8)
5 !$xmp align a(*,i) with t(i)

```

When an asterisk is specified as a subscript in a dimension of the target array in the align directive, the dimension is “collapsed” (i.e. not distributed). In the sample program above, the first dimension of the array *a* is distributed onto the node array *p* while the second dimension is collapsed.




Fig. 17 Example of array alignment (3).

In XMP/C, $a[0:2][:]$ will be allocated on $p[0]$ while, in XMP/Fortran, $a(:, 1:2)$ will be allocated on $p(1)$.

The programmer also can align an n -dimensional array with an m -dimensional template for $n < m$.

XscalableMP C

```

#pragma xmp nodes p[2][2]
#pragma xmp template t[8][8]
#pragma xmp distribute t[block][block] onto p
int a[8];
5 #pragma xmp align a[i] with t[i][*]

```

XscalableMP Fortran

```

!$xmp nodes p(2,2)
!$xmp template t(8,8)
!$xmp distribute t(block,block) onto p
integer :: a(8)
5 !$xmp align a(i) with t(*,i)

```

When an asterisk is specified as a subscript in a dimension of the target template in the align directive, the array will be “replicated” along the axis of the dimension.




Fig. 18 Example of array alignment (4).

In XMP/C, $a[0:4]$ will be replicated and allocated on $p[0][0]$ and $p[0][1]$ while, in XMP/Fortran, $a(1:4)$ will be allocated on $p(1, 1)$ and $p(2, 1)$.

2.5 Dynamic Allocation of Distributed Array

This section explains how distributed (i.e. global) arrays are allocated at runtime. The basic procedure is common in XMP/C and XMP/Fortran with a few specific difference.

```

XscalableMP C
#pragma xmp nodes p[4]
#pragma xmp template t[N]
#pragma xmp distribute t[block] onto p
float *a;
5 #pragma xmp align a[i] with t[i]
  :
a = xmp_malloc(xmp_desc_of(a), N);

```

In XMP/C, first, declare a pointer of the type of the target array; second, align it as if it were an array; finally, allocate memory for it with the `xmp_malloc()` function. `xmp_desc_of()` is an intrinsic/builtin function that returns the descriptor of the XMP object (i.e. nodes, templates, or global arrays) specified by the argument.

```

XscalableMP Fortran
!$xmp nodes p(4)
!$xmp template t(N)
!$xmp distribute t(block) onto p
real, allocatable :: a(:)
5 !$xmp align a(i) with t(i)

allocate(a(N))

```

In XMP/Fortran, first, declare an allocatable array; second, align it; finally, allocate memory for it with the `allocate` statement.

For multi-dimensional arrays, the procedure is the same as that for one-dimensional ones, as follows:

	XcalableMP C
	<pre>#pragma xmp nodes p[2][2] #pragma xmp template t[N1][N2] #pragma xmp distribute t[block][block] onto p float (*a)[N2]; 5 #pragma xmp align a[i][j] with t[i][j] : a = (float (*)(N2))xmp_malloc(xmp_desc_of(a), N1, N2);</pre>
	XcalableMP Fortran
	<pre>!\$xmp nodes p(2,2) !\$xmp template t(N2,N1) !\$xmp distribute t(block,block) onto p real, allocatable :: a(:, :) 5 !\$xmp align a(j,i) with t(j,i) : allocate(a(N2,N1))</pre>

Note: If the size of the template is not fixed until runtime, the programmer have to fix it at runtime with the `template_fix` construct.

2.6 template_fix Construct

The `template_fix` construct fixes the shape and/or the distribution of an unfixed template.

	XcalableMP C
	<pre>#pragma xmp nodes p[4] #pragma xmp template t[:] #pragma xmp distribute t[block] onto p double *a; 5 #pragma xmp align a[i] with t[i] int n = 100; #pragma xmp template_fix t[n] a = xmp_malloc(xmp_desc_of(a), n);</pre>

XcalableMP Fortran

```

!$xmp nodes p(4)
!$xmp template t(:)
!$xmp distribute t(block) onto p
real, allocatable :: a(:)
5 integer :: n
!$xmp align a(i) with t(i)

n = 100
!$xmp template_fix t(n)
10 allocate(a(n))

```

In the above sample code, first, a template `t` whose size is unfixed (“:”) is declared; second, a pointer `a`, in XMP/C, or an allocatable array `a`, in XMP/Fortran, is aligned with the template; third, the size of the template is fixed with a `template_fix` construct; finally, the pointer or the allocatable array is allocated with the `xmp_malloc()` builtin function in XMP/C or the `allocate` statement in XMP/Fortran, respectively.

Note: The `template_fix` constructs can be applied to a template only once.

This construct can also be used to fix a mapping array of a template that is distributed in “`gblock(*)`” at declaration.

XcalableMP C

```

#pragma xmp nodes p[4]
#pragma xmp template t[:]
#pragma xmp distribute t[gblock(*)] onto p
double *a;
5 #pragma xmp align a[i] with t[i]

int n = 100;
int m[] = {40,30,20,10};

10 #pragma xmp template_fix[gblock(m)] t[n]
a = xmp_malloc(xmp_desc_of(a), n);

```

XcalableMP Fortran

```

!$xmp nodes p(4)
!$xmp template t(:)
!$xmp distribute t(gblock) onto p
real, allocatable :: a(:)
5 integer :: n, m(4)
!$xmp align a(i) with t(i)

```



```

n = 100
m(:) = (/40,30,20,10/)
10 !$xmp template_fix(gblock(m)) t(n)
allocate(a(n))

```

3 Work Mapping

3.1 task and tasks Construct

The `task` construct defines a *task* that is executed by a specified node set. The `tasks` construct asserts that surrounding `task` constructs can be executed in parallel.

3.1.1 task Construct

When a node encounters a *task* construct at runtime, it executes the associated block (called a task) if it is included by the node set specified by the `on` clause; otherwise, it skips the execution of the block.

XcalableMP C

```

#include <stdio.h>
#pragma xmp nodes p[4]

int main(){
5   int num = xmpc_node_num();
  #pragma xmp task on p[1:3]
  {
    printf("%d: Hello\n", num);
  }
10
  return 0;
}

```

XcalableMP Fortran

```

program main
!$xmp nodes p(4)
  integer :: num

5   num = xmp_node_num()
!$xmp task on p(2:4)
  write(*,*) num, ": Hello"
!$xmp end task

10 end program main

```

In the above example, nodes $p[1]$, $p[2]$, and $p[3]$ invokes the `printf()` function, and $p[1]$ outputs “1: Hello” in XMP/C; $p(2)$, $p(3)$, and $p(4)$ execute the `write` statement, and $p(2)$ outputs “2: Hello” in XMP/Fortran.

Note that a new node set is generated by each task construct. Let’s consider inserting a `bcast` construct into the task.

XcalableMP C
<pre>#pragma xmp task on p[1:3] { #pragma xmp bcast (num) }</pre>

XcalableMP Fortran
<pre>!\$xmp task on p(2:4) !\$xmp bcast (num) !\$xmp end task</pre>

This `bcast` construct is executed by the node set specified by the `on` clause of the task construct. Thus, the node $p[1]$ broadcasts the value of `num` to $p[2]$ and $p[3]$ in XMP/C, and $p(2)$ to $p(3)$ and $p(4)$ in XMP/Fortran.

XMP/C

	p[0]	p[1]	p[2]	p[3]	
num =	0	1	2	3	
num =	0	1	1	1	} Region by task directive
		bcast directive			
num =	0	1	1	1	

XMP/Fortran

	p(1)	p(2)	p(3)	p(4)	
num =	1	2	3	4	
num =	1	2	2	2	} Region by task directive
		bcast directive			
num =	1	2	2	2	

Fig. 19 Example of task construct (1).

The `bcast` construct in the above code is equivalent to that in the following code, where it is executed by a new node set `q` that is explicitly declared.

XcalableMP C
<pre>#pragma xmp nodes q[3] = p[1:3] #pragma xmp bcast (num) on q</pre>
XcalableMP Fortran
<pre>!\$xmp nodes q(3) = p(2:4) !\$xmp bcast (num) on q</pre>

Note that the task is executed by the node set specified by the `on` clause. Therefore, `xmpc_node_num()` and `xmp_node_num()` return the id in the node set.

For example, consider inserting `xmpc_node_num()` or `xmp_node_num()` into the task in the first program.

XcalableMP C
<pre>#include <stdio.h> #pragma xmp nodes p[4] int main(){ 5 #pragma xmp task on p[1:3] { printf("%d: Hello\n", xmpc_node_num()); } 10 return 0; }</pre>
XcalableMP Fortran
<pre>program main !\$xmp nodes p(4) !\$xmp task on p(2:4) 5 write(*,*) xmp_node_num(), ": Hello" !\$xmp end task end program main</pre>

The node `p[1]` outputs “0: Hello” in XMP/C, and `p(2)` “1: Hello” in XMP/Fortran.

Note: A new node set should be collectively generated by all of the executing nodes at the point of a `task` construct unless it is surrounded by a `tasks` construct. Therefore, in the above example, `p[0]` in XMP/C and `p(1)` in XMP/Fortran must execute the `task` construct.

3.1.2 tasks Construct

Let's consider that each of two tasks invokes a function.

	XcalableMP C
	<code>#pragma xmp nodes p[4]</code>
	<code>#pragma xmp task on p[0:2]</code>
	<code>{</code>
5	<code>func_a();</code>
	<code>}</code>
	<code>#pragma xmp task on p[2:2]</code>
	<code>{</code>
	<code>func_b();</code>
10	<code>}</code>


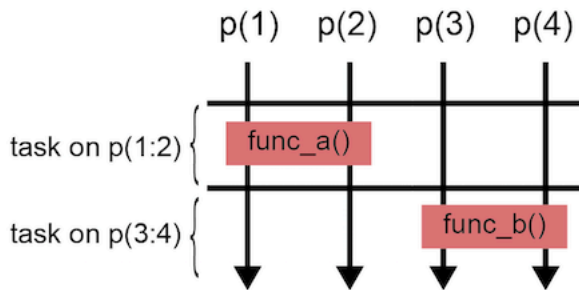
	XcalableMP Fortran
	<code>!\$xmp nodes p(4)</code>
	<code>!\$xmp task on p(1:2)</code>
	<code>call func_a()</code>
5	<code>!\$xmp end task</code>
	<code>!\$xmp task on p(3:4)</code>
	<code>call func_b()</code>
	<code>!\$xmp end task</code>

In the above example, the two tasks cannot be executed in parallel because those on clauses must be evaluated by all of the executing nodes.

In such a case, the programmer must specify a `tasks` construct surrounding the tasks is needed to execute them in parallel.

	XcalableMP C
	<code>#pragma xmp nodes p[4]</code>
	<code>#pragma xmp tasks</code>
	<code>{</code>
5	<code>#pragma xmp task on p[0:2]</code>
	<code>{</code>
	<code>func_a();</code>
	<code>}</code>
	<code>#pragma xmp task on p[2:2]</code>
10	<code>{</code>
	<code>func_b();</code>
	<code>}</code>
	<code>}</code>

	XcalableMP Fortran
	<code>!\$xmp nodes p(4)</code>

XMP/C**XMP/Fortran****Fig. 20** Example of task construct (2).

```



!$xmp tasks
!$xmp task on p(1:2)
5  call func_a()
!$xmp end task
!$xmp task on p(3:4)
   call func_b()
!$xmp end task
10 !$xmp end tasks

```

Because the node sets specified by the `on` clauses of the `task` constructs surrounded by a `tasks` construct are disjoint, they can be executed in parallel.

3.2 loop Construct

The `loop` construct is used to parallelize a loop.

XMP/C**XMP/Fortran****Fig. 21** Example of tasks construct.

XcalableMP C

```
#pragma xmp loop on t[i]
  for (int i = 0; i < 10; i++)
    a[i] = i;
```

XcalableMP Fortran

```
!$xmp loop on t(i)
  do i = 1, 10
    a(i) = i
  end do
```

The loop directive above specifies that the iteration i of the following loop is executed by the node that owns the template element $t[i]$ or $t(i)$, which is specified in the on clause.

Such a loop must satisfy the following two conditions:

1. There is no data/control dependence among the iterations. In other words, the iterations of the loop can be executed in any order to produce the same result.
2. Elements of distributed arrays, if any, are accessed only by the node(s) that own(s) the elements.

The programs below are examples of a right loop directive and a loop statement. The condition 1 is satisfied because *i* is the only one index of the distributed array *a* that is accessed within the loop, and the condition 2 is also satisfied because the indices of the template in the *on* clause of the loop directive is identical to that of the distributed array.

```

XcalableMP C
#pragma xmp nodes p[2]
#pragma xmp template t[10]
#pragma xmp distribute t[block] onto p

5 int main(){
    int a[10];
    #pragma xmp align a[i] with t[i]

    #pragma xmp loop on t[i]
10   for(int i=0;i<10;i++)
        a[i] = i;

    return 0;
}

```

```

XcalableMP Fortran
program main
!$xmp nodes p(2)
!$xmp template t(10)
!$xmp distribute t(block) onto p
5   integer a(10)
!$xmp align a(i) with t(i)

!$xmp loop on t(i)
    do i=1, 10
10      a(i) = i
    enddo

end program main

```

Then, is it possible to parallelize the loops in the below example where the loop bounds are shrunk from the above?

```

XcalableMP C
#pragma xmp loop on t[i]
    for(int i=1;i<9;i++)
        a[i] = i;

```

```

XcalableMP Fortran
!$xmp loop on t(i)
    do i=2, 9

```

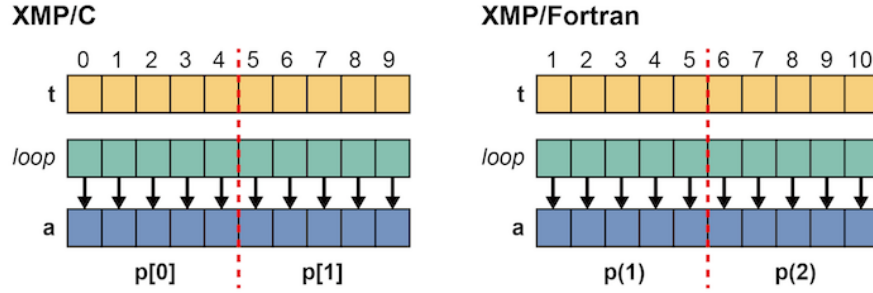


Fig. 22 Example of loop construct (1).

```

a(i) = i
enddo

```

In this case, the conditions 1 and 2 are satisfied and therefore it is possible to parallelize them. In XMP/C, $p[0]$ processes the indices from one to four and $p[1]$ from five to eight. In XMP/Fortran, $p(1)$ processes the indices from two to five and $p(2)$ from six to nine.

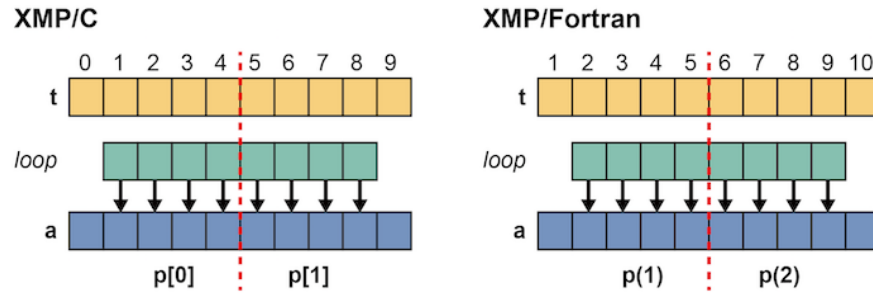


Fig. 23 Example of loop construct (2).

Next, is it possible to parallelize the below loops in which the index of the distributed array is different?

```

XcalableMP C
#pragma xmp loop on t[i]
for(int i=1;i<9;i++)
  a[i+1] = i;

XcalableMP Fortran
!$xmp loop on t(i)
do i=2, 9
  a(i+1) = i
enddo

```


In this case, the condition 1 is satisfied but 2 is not, and therefore it is not possible to parallelize them. In XMP/C, `p[0]` tries to access `a[5]` but does not own it. In XMP/Fortran, `p(1)` tries to access `a(6)` but does not own it.

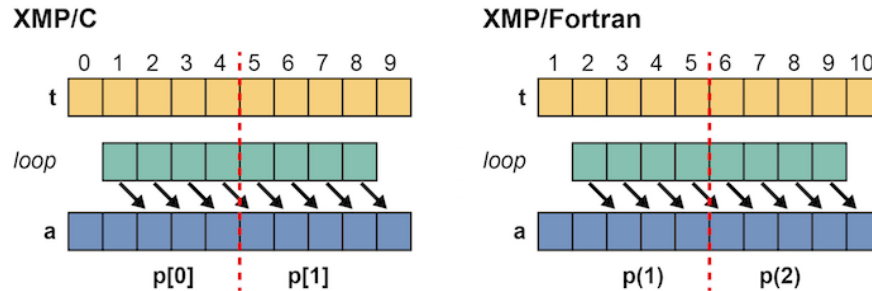


Fig. 24 Example of loop construct (3).

3.2.1 Reduction Computation

The serial programs below are examples of a reduction computation.

```

C
#include <stdio.h>

int main(){
    int a[10], sum = 0;
5   for(int i=0;i<10;i++){
        a[i] = i+1;
        sum += a[i];
    }
10  printf("%d\n", sum);

    return 0;
}

Fortran
program main
    integer :: a(10), sum = 0

    do i=1, 10
5       a(i) = i
        sum = sum + a(i)
    enddo

```

```

10  write(*,*) sum
    end program main

```

If the above loops are parallelized just by adding a loop directive, the value of the variable `sum` varies from node to node because it is calculated separately on each node. The value should be *reduced* to produce the right result.

XcalableMP C

```

5  #pragma xmp loop on t[i]
    for(int i=0;i<10;i++){
        a[i] = i+1;
        sum += a[i];
    }

```

XcalableMP Fortran

```

5  !$xmp loop on t(i)
    do i=1, 10
        a(i) = i
        sum = sum + a(i)
    enddo

```




Fig. 25 Example of reduction computation (1).

Then, to correct the error in the above code, add a **reduction** clause to the loop directive as follows.

XcalableMP C

```

5  #include <stdio.h>
    #pragma xmp nodes p[2]
    #pragma xmp template t[10]
    #pragma xmp distribute t[block] onto p
10 #pragma xmp loop on t[i] reduction(+:sum)

```

```

    for(int i=0;i<10;i++){
        a[i] = i+1;
        sum += a[i];
    }
15  printf("%d\n", sum);

    return 0;
}

```

XscalableMP Fortran

```

program main
!$xmp nodes p(2)
!$xmp template t(10)
!$xmp distribute t(block) onto p
5  integer :: a(10), sum = 0
!$xmp align a(i) with t(i)

!$xmp loop on t(i) reduction(+:sum)
do i=1, 10
10  a(i) = i
    sum = sum + a(i)
enddo

    write(*,*) sum
15
end program main

```

An operator and target variables for reduction computation are specified in a reduction clause. In the above examples, a “+” operator and a target variable sum are specified for the reduction computation to produce a total sum among nodes.

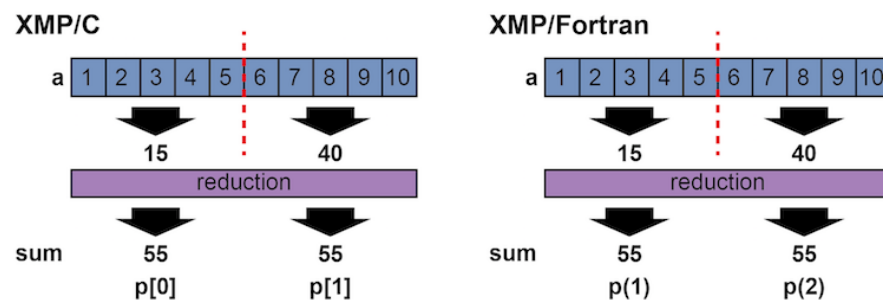


Fig. 26 Example of reduction computation (2).

Operations that can be specified as an operator in a `reduction` clause are limited to the following associative ones.

	C
	<code>+</code>
	<code>*</code>
	<code>-</code>
	<code>&</code>
5	<code> </code>
	<code>^</code>
	<code>&&</code>
	<code> </code>
	<code>max</code>
10	<code>min</code>
	<code>firstmax</code>
	<code>firstmin</code>
	<code>lastmax</code>
	<code>lastmin</code>

	Fortran
	<code>+</code>
	<code>*</code>
	<code>-</code>
	<code>.and.</code>
5	<code>.or.</code>
	<code>.eqv.</code>
	<code>.neqv.</code>
	<code>max</code>
	<code>min</code>
10	<code>iand</code>
	<code>ior</code>
	<code>ieor</code>
	<code>firstmax</code>
	<code>firstmin</code>
15	<code>lastmax</code>
	<code>lastmin</code>

Note: The total result is calculated by combining the partial results on all nodes. The ordering of the combination is unspecified. Hence, if the target variable is a type of floating point (e.g. `float` in XMP/C or `real` in XMP/Fortran), the difference of the order can make a little bit difference in the result value from that in the original serial execution.

3.2.2 Parallelizing Nested Loop

Parallelization of nested loops can be specified similarly to a single one, as follows.

	XcalableMP C
	<pre> #pragma xmp nodes p[2][2] #pragma xmp template t[10][10] #pragma xmp distribute t[block][block] onto p 5 int main(){ int a[10][10]; #pragma xmp align a[i][j] with t[i][j] #pragma xmp loop on t[i][j] 10 for(int i=0;i<10;i++) for(int j=0;j<10;j++) a[i][j] = i*10+j; return 0; 15 }</pre>
	XcalableMP Fortran
	<pre> program main !\$xmp nodes p(2,2) !\$xmp template t(10,10) !\$xmp distribute t(block,block) onto p 5 integer :: a(10,10) !\$xmp align a(j,i) with t(j,i) !\$xmp loop on t(j,i) do i=1, 10 10 do j=1, 10 a(j,i) = i*10+j enddo enddo 15 end program main</pre>

3.3 array Construct

The array construct is for work mapping of array assignment statements.

	XcalableMP C
	<pre> #pragma xmp align a[i] with t[i] :</pre>

```
#pragma xmp array on t[0:N]
a[0:N] = 1.0;
```

_____ XcalableMP Fortran _____

```
!$xmp align a(i) with t(i)
:
!$xmp array on t(1:N)
a(1:N) = 1.0
```

The above is equivalent to the below.

_____ XcalableMP C _____

```
#pragma xmp align a[i] with t[i]
:
#pragma xmp loop on t[i]
for(int i=0;i<N;i++)
5   a[i] = 1.0;
```

_____ XcalableMP Fortran _____

```
!$xmp align a(i) with t(i)
:
!$xmp loop on t(i)
do i=1, N
5   a(i) = 1.0
enddo
```

This construct can also be applied to multi-dimensional arrays.

_____ XcalableMP C _____

```
#pragma xmp align a[i][j] with t[i][j]
:
#pragma xmp array on t[:, :]
a[:, :] = 1.0;
```

_____ XcalableMP Fortran _____

```
!$xmp align a(j,i) with t(j,i)
:
!$xmp array on t(:, :)
a(:, :) = 1.0
```

Note: The template appearing in the on clause must have the same shape as the arrays in the following statement. The right-hand side value in this construct must be identical among all nodes because the array construct is a global (i.e. collective) operation.

4 Data Communication

4.1 shadow Directive and reflect Construct

Stencil computation frequently appears in scientific simulation programs, where, to update an array element $a[i]$, its neighboring elements $a[i-1]$ and $a[i+1]$ are referenced. If $a[i]$ is on the boundary region of a block-distributed array on a node, $a[i+1]$ may reside on another (neighboring) node.

Since it involves large overhead to copy $a[i+1]$ from the neighboring node to update each $a[i]$, a technique of copying collectively the elements on the neighboring node to the area added to the distributed array on each node is usually adopted. In XMP, such additional area is called “shadow.”

4.1.1 Declaring Shadow

Shadow areas can be declared with the `shadow` directive. In the example below, an array `a` has shadow areas of width one on both the lower and upper bounds.

	XcalableMP C
	<pre>#pragma xmp nodes p[4] #pragma xmp template t[16] #pragma xmp distribute t[block] onto p double a[16]; 5 #pragma xmp align a[i] with t[i] #pragma xmp shadow a[1]</pre>
	XcalableMP Fortran
	<pre>!\$xmp nodes p(4) !\$xmp template t(16) !\$xmp distribute t(block) onto p real :: a(16) 5 !\$xmp align a(i) with t(i) !\$xmp shadow a(1)</pre>



Fig. 27 Example of shadow directive (1).

In the figure above, colored elements are those that each node owns and white ones are shadow.

Note: Arrays distributed in a cyclic manner cannot have shadow.

In some programs, it is natural that the widths of the shadow area on the lower and upper bounds are different. There is also a case where the shadow area exists only on either of the bounds. In the example below, it is declared that a distributed array `a` has a shadow area of width one only on the upper bound.

	XcalableMP C
	<code>#pragma xmp nodes p[4]</code>
	<code>#pragma xmp template t[16]</code>
	<code>#pragma xmp distribute t(block) onto p</code>
	<code>double a[16];</code>
5	<code>#pragma xmp align a[i] with t[i]</code>
	<code>#pragma xmp shadow a[0:1]</code>

	XcalableMP Fortran
	<code>!\$xmp nodes p(4)</code>
	<code>!\$xmp template t(16)</code>
	<code>!\$xmp distribute t(block) onto p</code>
	<code>real :: a(16)</code>
5	<code>!\$xmp align a(i) with t(i)</code>
	<code>!\$xmp shadow a(0:1)</code>

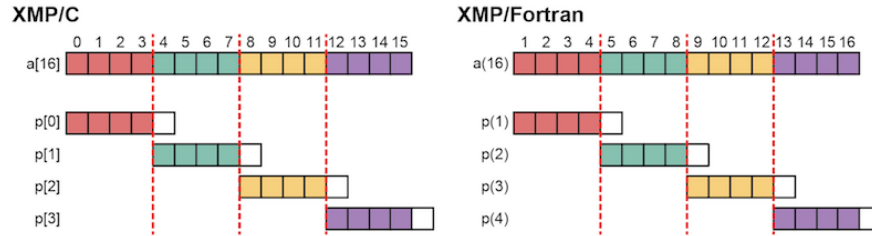


Fig. 28 Example of shadow directive (2).

The values on the left- and right-hand sides of a colon designate the widths on the lower and upper bounds, respectively.

4.1.2 Updating Shadow

To copy data to shadow areas from neighboring nodes, use the `reflect` construct. In the example below, the shadow areas of an array `a` that are of width one on both the upper and lower bounds are updated.

```

XcalableMP C
#pragma xmp reflect (a)

#pragma xmp loop on t[i]
for(int i=1;i<15;i++)
5   a[i] = (a[i-1] + a[i] + a[i+1])/3;

XcalableMP Fortran
!$xmp reflect (a)

!xmp loop on t(i)
do i=2, 15
5   a(i) = (a(i-1) + a(i) + a(i+1))/3
enddo

```

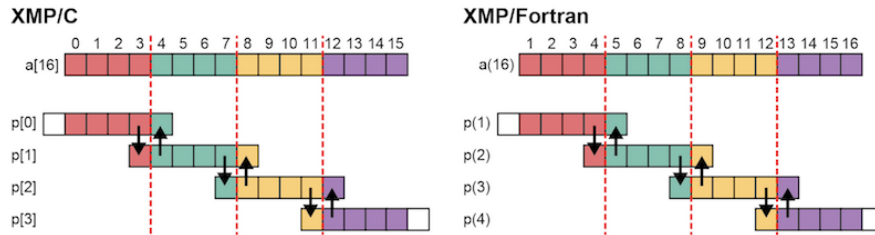


Fig. 29 Example of reflect construct (1).

With this `reflect` directive, in XMP/C, node `p[1]` sends an element `a[4]` to the shadow area on the upper bound on node `p[0]` and `a[7]` to the shadow area on the lower bound on `p[2]`; `p[0]` sends an element `a[3]` to the shadow area on the lower bound on `p[1]`, and `p[2]` sends `a[8]` to the shadow area on the upper bound on `p[1]`.

Similarly, in XMP/Fortran, node `p(2)` sends an element `a(5)` to the shadow area on the upper bound on node `p(1)` and `a(8)` to the shadow area on the lower bound on `p(3)`; `p(1)` sends an element `a(4)` to the shadow area on the lower bound on `p(2)`, and `p(3)` sends `a(9)` to the shadow area on the upper bound on `p(2)`.

The default behavior of a `reflect` directive is to update the whole of the shadow area declared by the `shadow` directive. However, there are some cases where a specific part of the shadow area is to be updated to reduce the communication cost at a point of the code.

To update only a specific part of the shadow area, add the `width` clause to the `reflect` directive.

The values on the left- and right-hand sides of a colon in the `width` clause designate the widths on the lower and upper bounds to be updated, respectively. In the example below, only the shadow area on the upper bound is updated.

```

XcalableMP C
#pragma xmp reflect (a) width(0:1)

XcalableMP Fortran
!$xmp reflect (a) width(0:1)

```

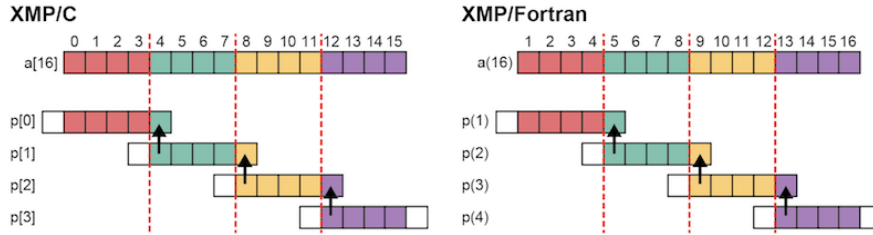


Fig. 30 Example of `reflect` construct (2).

Note: If the widths of the shadow areas to be updated on the upper and lower bounds are equal, that is, for example, `width(1:1)`, you can abbreviate it as `width(1)`.

Note: It is not possible to update the shadow area on a particular node because `reflect` is a collective operation.

The `reflect` directive does not update either the shadow area on the lower bound on the leading node or that on the upper bound on the last node. However, the values in such areas are needed for stencil computation if periodic boundary conditions are used in the computation.

To update such areas, add a `periodic` qualifier into the `width` clause. Let's look at the following example where an array having shadow areas of width one on both the lower and upper bounds appears.

```

XcalableMP C
#pragma xmp reflect (a) width(/periodic/1:1)

XcalableMP Fortran
!$xmp reflect (a) width(/periodic/1:1)

```

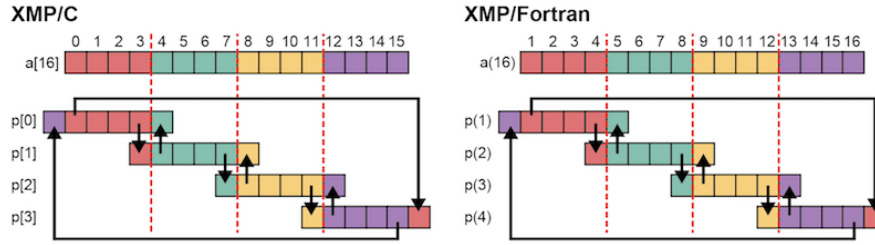


Fig. 31 Example of periodic reflect construct.

The periodic qualifier has the following effects, in addition to that of a normal reflect directive: in XMP/C, node $p[0]$ sends an element $a[0]$ to the shadow area on the upper bound on node $p[3]$, and $p[3]$ sends $a[15]$ to the shadow area on the lower bound on $p[0]$; in XMP/fortran, node $p(1)$ sends an element $a(1)$ to the shadow area on the upper bound on node $p(4)$, and $p(4)$ sends $a(16)$ to the shadow area on the lower bound on $p(1)$.

The shadow directive and reflect construct can be applied to arrays distributed in multiple dimensions. The following programs are the examples for two-dimensional distribution.

XcalableMP C

```
#pragma xmp nodes p[3][3]
#pragma xmp template t[9][9]
#pragma xmp distribute t[block][block] onto p
double a[9][9];
5 #pragma xmp align a[i][j] with t[i][j]
#pragma xmp shadow a[1][1]
:
#pragma xmp reflect (a)
```

XcalableMP Fortran

```
!$xmp nodes p(3,3)
!$xmp template t(9,9)
!$xmp distribute t(block,block) onto p
real :: a(9,9)
5 !$xmp align a(j,i) with t(j,i)
!$xmp shadow a(1,1)
:
!$xmp reflect (a)
```

The central node receives data from the surrounding eight nodes to update its shadow areas. The shadow areas of the other nodes are also updated, which is omitted in the figure.

For some applications, data from ordinal directions are not necessary. In such a case, the data communication from/to the ordinal directions can be avoided by adding the orthogonal clause to a reflect construct.

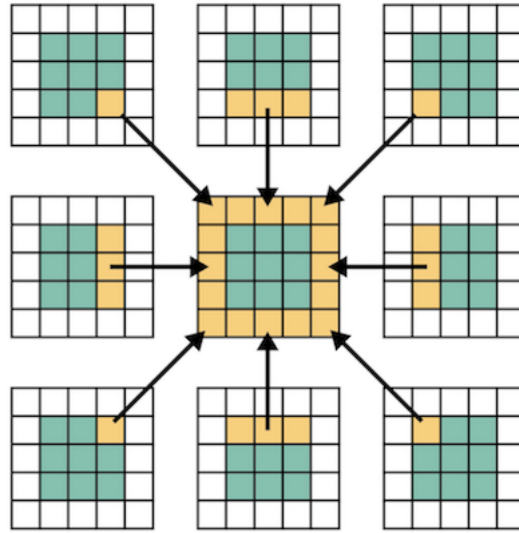


Fig. 32 Example of multi-dimensional shadow (1).

XcalableMP C

```
#pragma xmp reflect (a) orthogonal
```

XcalableMP Fortran

```
!$xmp reflect (a) orthogonal
```

Note: The `orthogonal` clause is effective only for arrays more than one dimension of which is distributed.

Besides, you can also add shadow areas to only specified dimension.

XcalableMP C

```
#pragma xmp nodes p[3]
#pragma xmp template t[9]
#pragma xmp distribute t[block] onto p
double a[9][9];
5 #pragma xmp align a[i][*] with t[i]
#pragma xmp shadow a[1][0]
:
#pragma xmp reflect (a)
```

XcalableMP Fortran

```
!$xmp nodes p[3]
!$xmp template t[9]
!$xmp distribute t[block] onto p
```

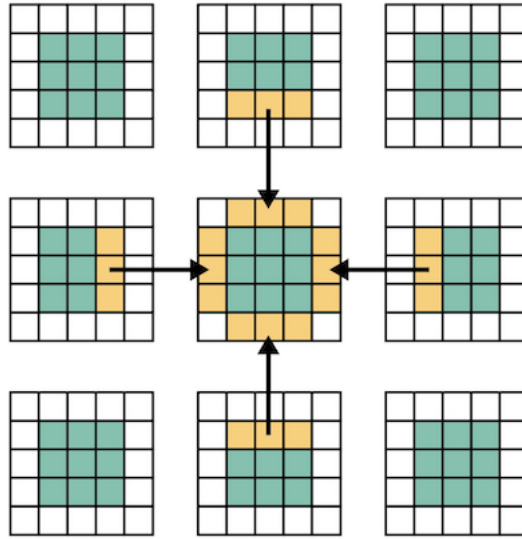


Fig. 33 Example of multi-dimensional shadow (2).

```

real :: a(9,9)
5 !$xmp align a(*,i) with t(i)
!$xmp shadow a(0,1)
:
!$xmp reflect (a)

```

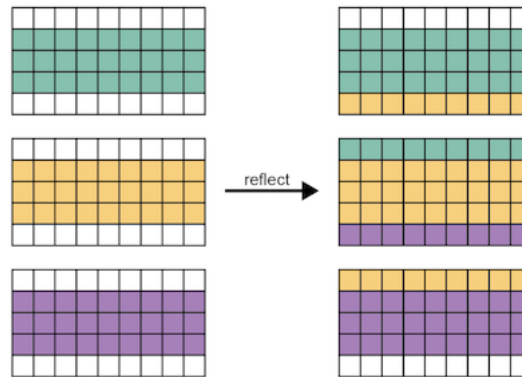


Fig. 34 Example of multi-dimensional shadow (3).

For the array `a`, 0 is specified as the shadow width in non-distributed dimensions.

4.2 gmove Construct

The programmers can specify a communication of distributed arrays in the form of assignment statements by using the `gmove` construct. In other words, with the `gmove` construct, any array assignment between two arrays (i.e. *global data movement*) that may involve inter-node communication can be specified.

There are three modes of `gmove`; “collective mode,” “in mode,” and “out mode.”

4.2.1 Collective Mode

The global data movement involved by a *collective* `gmove` is performed collectively, and results in implicit synchronization among the executing nodes.

	XcalableMP C
	<pre> #pragma xmp nodes p[4] #pragma xmp template t[16] #pragma xmp distribute t[block] onto p int a[16], b[16]; 5 #pragma xmp align a[i] with t[i] #pragma xmp align b[i] with t[i] : #pragma xmp gmove a[9:5] = b[0:5]; </pre>
	XcalableMP Fortran
	<pre> !\$xmp nodes p(4) !\$xmp template t(16) !\$xmp distribute t(block) onto p integer :: a(16), b(16) 5 !\$xmp align a(i) with t(i) !\$xmp align b(i) with t(i) : !\$xmp gmove a(10:14) = b(1:5) </pre>




Fig. 35 Collective `gmove` (1).

In XMP/C, $p[0]$ sends $b[0]$ - $b[3]$ to $p[2]$ - $p[3]$, and $p[1]$ sends $b[4]$ to $p[3]$. Similarly, in XMP/Fortran, $p(1)$ sends $b(1)$ - $b(4)$ to $p(3)$ - $p(4)$, and $p(2)$ sends $b(5)$ to $p(4)$.

```

XcalableMP C
#pragma xmp nodes p[4]
#pragma xmp template t1[16]
#pragma xmp template t2[16]
#pragma xmp distribute t1[cyclic] onto p
5 #pragma xmp distribute t2[block] onto p
  int a[16], b[16];
  #pragma xmp align a[i] with t1[i]
  #pragma xmp align b[i] with t2[i]
  :
10 #pragma xmp gmove
    a[9:5] = b[0:5];

XcalableMP Fortran
!$xmp nodes p(4)
!$xmp template t1(16)
!$xmp template t2(16)
!$xmp distribute t1(cyclic) onto p
5 !$xmp distribute t2(block) onto p
  integer :: a(16), b(16)
  !$xmp align a(i) with t1(i)
  !$xmp align b(i) with t2(i)
  :
10 !$xmp gmove
    a(10:14) = b(1:5)

```

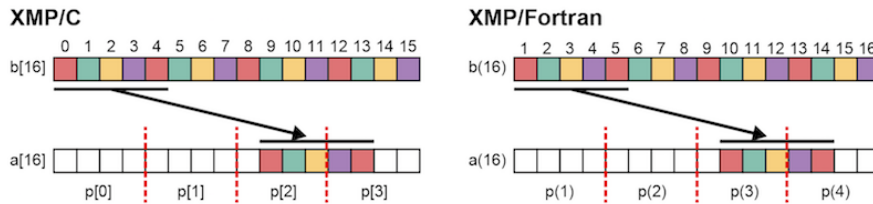


Fig. 36 Collective gmove (2).

While array a is distributed in a cyclic manner, array b is distributed in a block manner.

In XMP/C, $p[0]$ sends $b[0]$ and $b[4]$ to $p[2]$ and $p[3]$. $p[1]$ sends $b[1]$ to $p[2]$. Each element of $p[2]$ and $p[3]$ will be copied locally. Similarly, in XMP/Fortran, $p(1)$ sends $b(1)$ and $b(5)$ to $p(3)$ and $p(4)$. $p(2)$ sends $b(2)$ to $p(3)$. Each element of $p(3)$ and $p(4)$ will be copied locally.

By using this method, the distribution of an array can be “changed” during computation.

```

XcalableMP C
#pragma xmp nodes p[4]
#pragma xmp template t1[16]
#pragma xmp template t2[16]
int W[4] = {2,4,8,2};
5 #pragma xmp distribute t1[gblock(W)] onto p
  #pragma xmp distribute t2[block] onto p
  int a[16], b[16];
  #pragma xmp align a[i] with t1[i]
  #pragma xmp align b[i] with t2[i]
10 :
  #pragma xmp gmove
    a[:] = b[:];

XcalableMP Fortran
!$xmp nodes p(4)
!$xmp template t1(16)
!$xmp template t2(16)
integer :: W(4) = (/2,4,7,3/)
5 !$xmp distribute t1(gblock(W)) onto p
  !$xmp distribute t2(block) onto p
  integer :: a(16), b(16)
  !$xmp align a(i) with t1(i)
  !$xmp align b(i) with t2(i)
10 :
  !$xmp gmove
    a(:) = b(:)

```

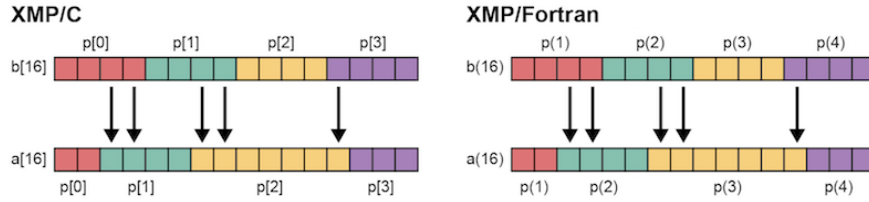


Fig. 37 Collective gmove (3).

In this example, the elements of an array *b* that is distributed in a block manner are copied to the corresponding elements of an array *a* that is distributed in a generalized-block manner. For the arrays *a* and *b*, communication occurs if corresponding elements reside in different nodes (arrows illustrate communication between nodes in the figures).

In the assignment statement, if a scalar (i.e. one element of an array or a variable) is specified on the right-hand side and an array section are specified on the left-hand side, a broadcast communication occurs for it.

```

XscalableMP C
#pragma xmp nodes p[4]
#pragma xmp template t[16]
#pragma xmp distribute t[block] onto p
int a[16], b[16];
5 #pragma xmp align a[i] with t[i]
  #pragma xmp align b[i] with t[i]
  :
#pragma xmp gmove
  a[9:5] = b[0];

XscalableMP Fortran
!$xmp nodes p(4)
!$xmp template t(16)
!$xmp distribute t(block) onto p
integer :: a(16), b(16)
5 !$xmp align a(i) with t(i)
  !$xmp align b(i) with t(i)
  :
!$xmp gmove
  a(10:14) = b(1)

```




Fig. 38 Collective gmove (4).

In this example, in XMP/C, an array element `b[0]` of node `p[0]` will be broadcasted to the specified array section on node `p[2]` and `p[3]`. Similarly, in XMP/Fortran, an array element `b(1)` of node `p(1)` will be broadcasted to the specified array section on node `p(3)` and `p(4)`.

Not only distributed arrays but also replicated arrays can be specified on the right-hand side.

```

XscalableMP C
#pragma xmp nodes p[4]
#pragma xmp template t[16]
#pragma xmp distribute t[block] onto p

```

```

int a[16], b[16], c;
5  #pragma xmp align a[i] with t[i]
    :
#pragma xmp gmove
    a[9:5] = b[0:5];

```

XscalableMP Fortran

```

!$xmp nodes p(4)
!$xmp template t(16)
!$xmp distribute t(block) onto p
integer :: a(16), b(16), c
5  !$xmp align a(i) with t(i)
    :
!$xmp gmove
    a(10:14) = b(1:5)

```

In this example, a replicated array *b* is locally copied to distributed array *a* without communication.

XscalableMP C

```

#pragma xmp nodes p[4]
#pragma xmp template t1[8]
#pragma xmp template t2[16]
#pragma xmp distribute t1[block] onto p
5  #pragma xmp distribute t2[block] onto p
int a[8][16], b[8][16];
#pragma xmp align a[i][*] with t1[i]
#pragma xmp align b[*][i] with t2[i]
    :
10 #pragma xmp gmove
    a[0][:] = b[0][:];

```



XscalableMP Fortran

```

!$xmp nodes p(4)
!$xmp template t1(8)
!$xmp template t2(16)
!$xmp distribute t1(block) onto p
5  !$xmp distribute t2(block) onto p
integer :: a(16,8), b(8,16)
!$xmp align a(*,i) with t1(i)
!$xmp align b(i,*) with t2(i)
    :
10 #pragma xmp gmove
    a(:,1) = b(:,1)

```

In this example, in XMP/C, $b[0][0:2]$ on $p[0]$, $b[0][2:2]$ of $p[1]$, $b[0][4:2]$ on $p[2]$ and $b[0][6:2]$ on $p[3]$ are copied to $a[0][:]$ on $p[0]$. Similarly, in XMP/Fortran, $b(1:2,1)$ on $p(1)$, $b(3:4,1)$ of $p(2)$, $b(5:6,1)$ on $p(3)$ and $b(7:8,1)$ on $p(4)$ are copied to $a(:,1)$ on $p(1)$.

XMP/C**XMP/Fortran****Fig. 39** Collective gmove (4).**4.2.2 In Mode**

The right-hand side data of the assignment, all or part of which may reside outside the executing node set, can be transferred from its owner nodes to the executing nodes with an *in* gmove.

	XcalableMP C
	<pre> #pragma xmp nodes p[4] #pragma xmp template t[4] #pragma xmp distribute t[block] onto p double a[4], b[4]; 5 #pragma xmp align a[i] with t[i] #pragma xmp align b[i] with t[i] : #pragma xmp task on p[0:2] #pragma xmp gmove in 10 a[0:2] = b[2:2] #pragma xmp end task </pre>
	XcalableMP Fortran
	<pre> !\$xmp nodes p(4) !\$xmp template t(4) !\$xmp distribute t(block) onto p real :: a(4), b(4) </pre>


```

5  !$xmp align a(i) with t(i)
   !$xmp align b(i) with t(i)
   :
   !$xmp task on p(1:2)
   !$xmp gmove in
10  a(1:2) = b(3:4)
   !$xmp end task

```

In this example, the `task` directive divides four nodes into two sets, the first-half and the second-half. A `gmove` construct that is in an *in* mode copies data using a *get* operation from the second-half node to the first-half node.

XMP/C



XMP/Fortran




Fig. 40 In `gmove`.

4.2.3 Out Mode

For the left-hand side data of the assignment, all or part of which may reside outside the executing node set, the corresponding elements can be transferred from the executing nodes to its owner nodes with an *out gmove* construct.

XcalableMP C
<pre> #pragma xmp nodes p[4] #pragma xmp template t[4] #pragma xmp distribute t[block] onto p double a[4], b[4]; 5 #pragma xmp align a[i] with t[i] #pragma xmp align b[i] with t[i] : #pragma xmp task on p[0:2] #pragma xmp gmove out 10 b[2:2] = a[0:2] #pragma xmp end task </pre>
XcalableMP Fortran
<pre> !\$xmp nodes p(4) !\$xmp template t(4) !\$xmp distribute t(block) onto p real :: a(4), b(4) 5 !\$xmp align a(i) with t(i) !\$xmp align b(i) with t(i) : !\$xmp task on p(1:2) !\$xmp gmove out 10 b(3:4) = a(1:2) !\$xmp end task </pre>



A gmove construct that is in *out* mode copies data using a *put* communication from the first-half nodes to the second-half nodes.

4.3 barrier Construct

The barrier construct executes a barrier synchronization.

XcalableMP C
<pre> #pragma xmp barrier </pre>
XcalableMP Fortran
<pre> !\$xmp barrier </pre>

You can specify a node set on which the barrier synchroniation is to be performed by using the *on* clause. In the below example, a barrier synchronization is performed among the first two nodes of *p*.

XMP/C**XMP/Fortran****Fig. 41** Out gmove.

XcalableMP C
#pragma xmp barrier on p[0:2]
XcalableMP Fortran
!\$xmp barrier on p(1:2)

4.4 reduction Construct

This construct performs a *reduction* operation. It has the same meaning as the *reduction* clause of the *loop* construct, but this construct can be specified anywhere as an *executable* construct.

XcalableMP C
#pragma xmp nodes p[4] : sum = xmpc_node_num() + 1; #pragma xmp reduction (+:sum)

```

XcalableMP Fortran
!$xmp nodes p(4)
:
sum = xmp_node_num()
!$xmp reduction (+:sum)

```

XMP/C

	p[0]	p[1]	p[2]	p[3]	
sum =	1	2	3	4	reduction(+:sum)
	<hr/>				
sum =	10	10	10	10	

XMP/Fortran

	p(1)	p(2)	p(3)	p(4)	
sum =	1	2	3	4	reduction(+:sum)
	<hr/>				
sum =	10	10	10	10	

Fig. 42 reduction construct (1).

You can specify the executing node set by using the `on` clause. In the below example, only the values on the last two of the four nodes are targeted by the reduction construct.

```

XcalableMP C
#pragma xmp nodes p[4]
:
sum = xmpc_node_num() + 1;
#pragma xmp reduction (+:sum) on p[2:2]

```

```

XcalableMP Fortran
!$xmp nodes p(4)
:
sum = xmp_node_num()
!$xmp reduction (+:sum) on p(3:4)

```

The operators you can use in the reduction construct are as follows:

```

XcalableMP C
+
*
-

```

XMP/C

	p[0]	p[1]	p[2]	p[3]	
sum =	1	2	3	4	
	<hr/>				
sum =	1	2	7	7	reduction (+:sum) on p[2:2]

XMP/Fortran

	p(1)	p(2)	p(3)	p(4)	
sum =	1	2	3	4	
	<hr/>				
sum =	1	2	7	7	reduction (+:sum) on p(3:4)

Fig. 43 reduction construct (2).

	&	
5		
	^	
	&&	
	max	
10	min	
XscalableMP Fortran		
	+	
	*	
	-	
	.and.	
5	.or.	
	.eqv.	
	.neqv.	
	max	
	min	
10	iand	
	ior	
	ieor	

Note: In contrast to the reduction clause of the loop construct, which precedes loops, the reduction construct does not accept operators of `firstmax`, `firstmin`, `lastmax`, and `lastmin`.

Note: Similar to the `reduction` clause, the `reduction` construct may generate slightly different results in a parallel execution from those in a sequential execution, because the results depends on the order of combining the value.

4.5 bcast Construct

The `bcast` construct broadcasts the values of the variables on the node specified by the `from` clause, that is, the *root node*, to the node set specified by the `on` clause. If there is no `from` clause, the first node of the executing node set is selected as the root node. If there is no `on` clause, the current executing node set of the construct is selected as the executing node set.

In the below example, the first node of the node set `p`, that is, `p[0]` or `p(1)`, is the root node.

XcalableMP C
<pre>#pragma xmp nodes p[4] : num = xmpc_node_num() + 1; #pragma xmp bcast (num)</pre>
XcalableMP Fortran
<pre>!\$xmp nodes p(4) : num = xmp_node_num() !\$xmp bcast (num)</pre>

In the below example, the last node, that is, `p[3]` or `p(4)`, is the `from` clause.

XcalableMP C
<pre>#pragma xmp nodes p[4] : num = xmpc_node_num() + 1; #pragma xmp bcast (num) from p[3]</pre>
XcalableMP Fortran
<pre>!\$xmp nodes p(4) : num = xmp_node_num() !\$xmp bcast (num) from p(4)</pre>

In the below example, only the last three of four nodes are included by the executing node set of the `bcast` construct.

XMP/C

	p[0]	p[1]	p[2]	p[3]	
num =	1	2	3	4	bcast (num)
num =	1	1	1	1	

XMP/Fortran

	p(1)	p(2)	p(3)	p(4)	
num =	1	2	3	4	bcast (num)
num =	1	1	1	1	

Fig. 44 bcast construct (1).**XMP/C**

	p[0]	p[1]	p[2]	p[3]	
num =	1	2	3	4	bcast (num) from p[3]
num =	4	4	4	4	

XMP/Fortran

	p(1)	p(2)	p(3)	p(4)	
num =	1	2	3	4	bcast (num) from p(4)
num =	4	4	4	4	

Fig. 45 bcast construct (2).

```

XcalableMP C
#pragma xmp nodes p[4]
:
sum = xmpc_node_num() + 1;
#pragma xmp bcast (sum) from p[3] on p[1:3]

```

XcalableMP Fortran

```
!$xmp nodes p(4)
:
sum = xmp_node_num()
!$xmp bcast (num) from p(4) on p(2:4)
```

XMP/C

	p[0]	p[1]	p[2]	p[3]	
num =	1	2	3	4	
num =	1	4	4	4	bcast (num) from p[3] on p[1:3]

XMP/Fortran

	p(1)	p(2)	p(3)	p(4)	
num =	1	2	3	4	
num =	1	4	4	4	bcast (num) from p(4) on p(2:4)

Fig. 46 bcast construct (3).

4.6 wait_async Construct

Communication directives (i.e. `reflect`, `gmove`, `reduction`, `bcast`, and `reduce_shadow`) can perform asynchronous communication if the `async` clause is added. The `wait_async` construct is used to guarantee the completion of such an asynchronous communication.

XcalableMP C

```
#pragma xmp bcast (num) async(1)
:
#pragma xmp wait_async (1)
```

XcalableMP Fortran

```
!$xmp bcast (num) async(1)
:
!$xmp wait_async (1)
```

Since the `bcast` directive has an `async` clause, communication may not be completed immediately after the `bcast` directive. The completion of that communication is guaranteed with the `wait_async` construct having the same value as that of the `async` clause. Therefore, between the `bcast` construct and the `wait_async` constructs, you may not reference the target variable of the `bcast` directive.

Hint: Asynchronous communication can be overlapped with the following computation to hide its overhead.

Note: Expressions that can be specified as *tags* in the `async` clause are of type `int`, in XMP/C, or `integer`, in XMP/Fortran.

4.7 reduce_shadow Construct

The `reduce_shadow` directive adds the value of a shadow object to the corresponding data object of the array.

	XcalableMP C
<div style="margin-left: 20px;">5</div> <div style="margin-left: 20px;">10</div>	<pre> #pragma xmp nodes p[2] #pragma xmp template t[8] #pragma xmp distribute t[block] onto p int a[8]; #pragma xmp align a[i] with t[i] #pragma xmp shadow a[1] : #pragma xmp loop on t[i] for(int i=0;i<8;i++) a[i] = i+1; #pragma xmp reflect (a) #pragma xmp reduce_shadow (a) </pre>
	XcalableMP Fortran
<div style="margin-left: 20px;">5</div>	<pre> !\$xmp nodes p(2) !\$xmp template t(8) !\$xmp distribute t(block) onto p integer a(8) !\$xmp align a(i) with t(i) !\$xmp shadow a(1) </pre>

```

10 | !$xmp loop on t(i)
    |   do i=1, 8
    |     a(i) = i
    |   enddo
    |
    | !$xmp reflect (a)
    | !$xmp reduce_shadow (a)

```

For the above example, in XMP/C, $a[3]$ on $p[0]$ has a value of eight, and $a[4]$ on $p[1]$ has a value of ten. Similarly, in XMP/Fortran, $a(4)$ of $p(1)$ has a value of eight, and $a(5)$ on $p(2)$ has a value of ten.

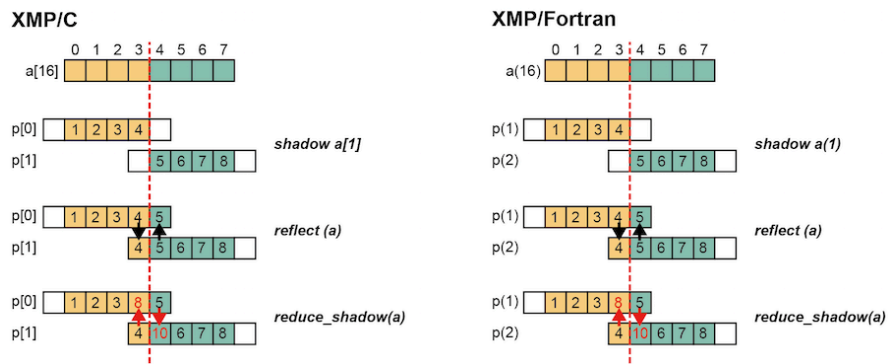


Fig. 47 `reduce_shadow` construct (1).

The programmers can add the `periodic` modifier to the width clause to reduce shadow objects to the corresponding data object periodically.

```

XscalableMP C
#pragma xmp reflect (a) width(/periodic/1)
#pragma xmp reduce_shadow (a) width(/periodic/1)

```

```

XscalableMP Fortran
!$xmp reflect (a) width(/periodic/1)
!$xmp reduce_shadow (a) width(/periodic/1)

```

In addition to the first example, in XMP/C, $a[0]$ on $p[0]$ has a value of two, and $a[7]$ on $p[1]$ has a value of 16. Similarly, in XMP/Fortran, $a(1)$ in $p(1)$ has a value of two, and $a(8)$ in $p(2)$ has a value of 16.

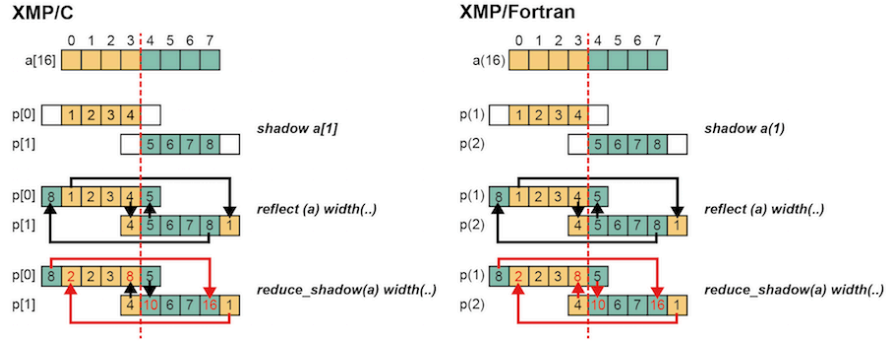


Fig. 48 reduce_shadow construct (2).

5 Local-view Programming

5.1 Introduction

The programmer can use coarrays to specify one-sided communication in the local-view model.

Depending on the environment, such one-sided communication might achieve better performance than global communication in the global-view model. However, it is more difficult and complicated to write parallel programs in the local-view model because the programmer must specify every detail of parallelization, such as data mapping, work mapping, and communication.

The coarray feature in XMP/Fortran is upward-compatible with that in Fortran 2008; that in XMP/C is defined as an extension to the base language.

An execution entity in local-view XMP programs is referred to as an “image” while a “node” in global-view ones. These two words have almost the same meaning in XMP.

5.2 Coarray Declaration

XcalableMP C	
<code>int a[10]:[*];</code>	
XcalableMP Fortran	
<code>integer a(10)[*]</code>	

In XMP/C, the programmer declares a coarray by adding “:[*]” after the array declaration. In XMP/Fortran, the programmer declares a coarray by adding “[*]” after the array declaration.

Note: Based on Fortran 2008, coarrays should have the same size among all images.

Coarrays can be accessed in expressions by remote images as well the local image.

5.3 Put Communication

When a coarray appears in the left-hand side of an assignment statement, it involves *put* communication.

XcalableMP C
<pre>int a[10]:[*], b[10]; if (xmpc_this_image() == 0) a[0:3]:[1] = b[3:3];</pre>

XcalableMP Fortran
<pre>integer a(10)[*] integer b(10) if (this_image() == 1) then a(1:3)[2] = b(3:5) end if</pre>

The integer in the square bracket specifies the target image index. The image index is zero-based, in XMP/C, or one-based, in XMP/Fortran. `xmpc_this_image()` in XMP/C and `this_image()` in XMP/Fortran return the current image index.



In the above example, in XMP/C, an image zero puts `b[3:3]` to `a[0:3]` on image one; in XMP/Fortran, an image one puts `b(3:5)` to `a(1:3)` on image two. The following figure illustrates the put communication performed in the example.

5.4 Get Communication

When a coarray appears in the right-hand side of an assignment statement, it involves *get* communication.

XcalableMP C
<pre>int a[10]:[*], b[10]; if (xmpc_this_image() == 0) b[3:3] = a[0:3]:[1];</pre>

XcalableMP Fortran
<pre>integer a(10)[*] integer b(10)</pre>



XMP/C**XMP/Fortran****Fig. 49** Remote write to a coarray

```

if (this_image() == 1) then
  b(3:5) = a(1:3)[2]
end if

```

In the above example, in XMP/C, an image 0 gets $a[0:3]$ from an image 1 and copies it to $b[3:3]$; in XMP/Fortran, an image 1 gets $a(1:3)$ from an image 2 and copies it to $b(3:5)$ of an image 1. The following figure illustrates the get communication performed in the example.

XMP/C**XMP/Fortran****Fig. 50** Remote read from a coarray

Hint: As illustrated above, get communication involves an extra step to send a request to the target node. Put communication achieves better performance than get because there is no such extra step.

5.5 Synchronization

5.5.1 sync all

XcalableMP C

```
void xmp_sync_all(int *status)
```

XcalableMP Fortran

```
sync all
```

At “sync all,” each image waits until all issued one-sided communication is complete and then performs barrier synchronization among the all images.

In the above example, the left image puts data to the right image and both nodes invoke sync all. When both nodes returns from it, the execution continues to the following statements.

5.5.2 sync images

XcalableMP C

```
void xmp_sync_images(int num, int *image-set, int *status)
```

XcalableMP Fortran

```
sync images (image-set)
```

Each image in the specified image set waits until all one-sided communication issued is complete, and performs barrier synchronization among the images.

XcalableMP C

```
int image_set[3] = {0,1,2};
xmp_sync_images(3, image_set, NULL);
```

XcalableMP Fortran

```
integer :: image_set(3) = (/ 1, 2, 3/)
sync images (image_set)
```

5.5.3 sync memory

XcalableMP C

```
void xmp_sync_memory(int *status)
```

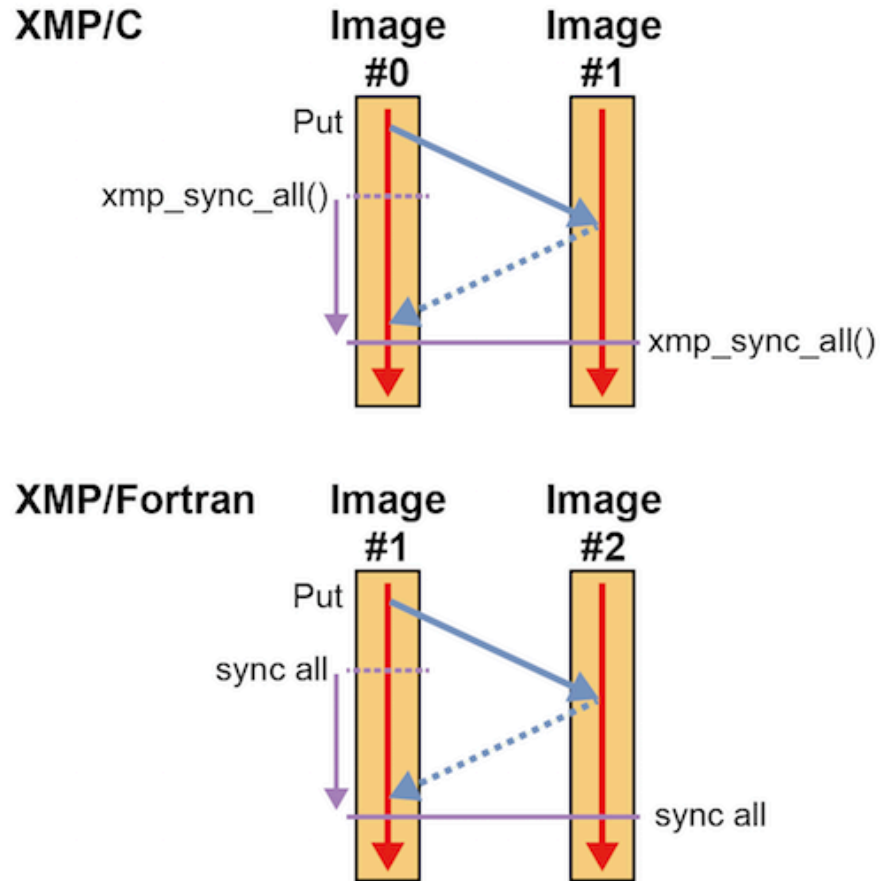


Fig. 51 sync all

XscalableMP Fortran

sync memory

Each image waits until all one-sided communication is complete. This function/statement does not imply barrier synchronization, unlike `sync all` and `sync images`, and therefore can be locally executed.

6 Procedure Interface

Procedure calls in XMP are almost the same as those in the base language. Procedure calls between other languages or to external libraries are also allowed if the base language supports them.

In the below example, a function/subroutine `sub1()` calls another function/subroutine `sub2()` with a distributed array `x` as an argument.

```

XcalableMP C
void sub1(){
  #pragma xmp nodes p[2]
  #pragma xmp template t[10]
  #pragma xmp distribute t[block] onto p
5   double x[10];
  #pragma xmp align x[i] with t[i]
    sub2(x);
}

10 void sub2(double a[10]){
  #pragma xmp nodes p[2]
  #pragma xmp template t[10]
  #pragma xmp distribute t[block] onto p
    double a[10];
15  #pragma xmp align a[i] with t[i]
    :
  }

XcalableMP Fortran
subroutine sub1()
  !$xmp nodes p(2)
  !$xmp template t(10)
  !$xmp distribute t(block) onto p
5   real x(10)
  !$xmp align x(i) with t(i)
    call sub2(x)
end subroutine

10 subroutine sub2(a)
  !$xmp nodes p(2)
  !$xmp template t(10)
  !$xmp distribute t(block) onto p
    real a(10)
15  !$xmp align a(i) with t(i)
    :
end subroutine

```

To handle a parameter or dummy argument as a global data in the callee procedure, the programmer need to explicitly distribute it with an `align` directive.

If no `align` directive is specified in the callee procedure for a parameter or dummy argument that is declared as a global data in the caller procedure, it is handled as if it were declared in the callee procedure as a local data on each node, as follows.

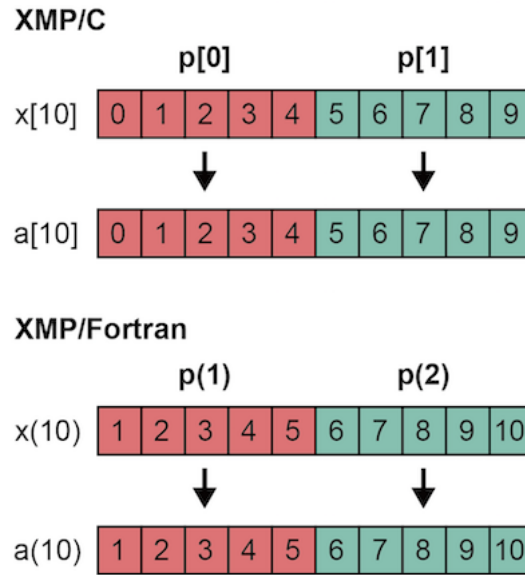


Fig. 52 Passing a global argument to a global parameter.

```

XcalableMP C
void sub1(){
  #pragma xmp nodes p[2]
  #pragma xmp template t[10]
  #pragma xmp distribute t[block] onto p
5   double x[10];
  #pragma xmp align x[i] with t[i]
  sub2(x);
}

10 void sub2(double a[5]){
  :
}

XcalableMP Fortran
subroutine sub1()
  !$xmp nodes p(2)
  !$xmp template t(10)
  !$xmp distribute t(block) onto p
5   real x(10)
  !$xmp align x(i) with t(i)
  call sub2(x)
end subroutine

```

```

10 subroutine sub2(a)
    real a(5)
    :
end subroutine

```



XMP/C**XMP/Fortran**

Fig. 53 Passing a global argument to a local parameter.

References

1. Robert W. Numrich and John Reid, “Co-Array Fortran for parallel programming”, ACM SIGPLAN Fortran Forum, Vol. 17, No. 2 (1998).
2. UPC Consortium, “UPC Specifications, v1.2”, Lawrence Berkeley National Lab (LBNL-59208) (2005).
3. David Callahan, Bradford L. Chamberlain and Hans P. Zima, “The Cascade High Productivity Language”, Proc. 9th Int’l. Workshop on High-Level Parallel Programming Models and Supportive Environments (HIPS 2004), pp. 52–60 (2004).