

XcalableMP Programming Model and Language

Hitoshi Murai, Masahiro Nakao, and Mitsuhsa Sato

Abstract XcalableMP (XMP) is a directive-based language extension of Fortran and C for distributed-memory parallel computers, and can be classified as a partitioned global address space (PGAS) language. One of remarkable characteristics of XMP is that it supports both of global- and local-view parallel programming. This chapter describes the programming model and language specification of XMP.

1 Introduction

Distributed-memory systems are generally used for large-scale simulations. To program such systems, Message Passing Interface (MPI) is widely adopted. However, programming with MPI is difficult because programmers must describe inter-process communications with consideration of the execution flow of the programs, which might cause deadlocks or wrong results.

To address this issue, a parallel language named High Performance Fortran (HPF) was proposed in 1991. With HPF, users can execute their serial programs in parallel by inserting minimal directives into them. If users specify data distribution with HPF directives, compilers do all other tasks for parallelization (e.g. communication generation and work distribution). However, HPF was not widely accepted eventually because the compilers' automatic processing prevents users from performance tuning

Hitoshi Murai
RIKEN Center for Computational Science, 7-1-26 Minatojima-minami-machi, Chuo-ku, Kobe,
Hyogo 650-0047, Japan, e-mail: h-murai@riken.jp

Masahiro Nakao
RIKEN Center for Computational Science, 7-1-26 Minatojima-minami-machi, Chuo-ku, Kobe,
Hyogo 650-0047, Japan, e-mail: masahiro.nakao@riken.jp

Mitsuhsa Sato
RIKEN Center for Computational Science, 7-1-26 Minatojima-minami-machi, Chuo-ku, Kobe,
Hyogo 650-0047, Japan, e-mail: msato@riken.jp

and the performance depends heavily on the environment (e.g. compiler version and hardware)

Note: For more detail, please refer: Ken Kennedy, Charles Koelbel and Hans Zima: The Rise and Fall of High Performance Fortran: An Historical Object Lesson, Proc. 3rd ACM SIGPLAN History of Programming Languages Conf. (HOPL-III), pp. 7-17-22 (2007).

In such circumstance, to develop a new parallel programming model that enables easy parallelization of existing serial programs and design a new language based on it, “the XMP Specification Working Group” was established in 2008. This group utilized the lessons from the experience of HPF to define a parallel language XcalableMP (XMP). The group was reorganized to one of the working groups of PC Cluster Consortium in 2011.

It is learned from the lessons of HPF that more automatic processing of compilers increases the gap between a program and its execution, and, as a result, decreases the usability of the language.

In XMP, users specify explicitly the details of parallel programs on the basis of compiler directives to make their execution easy-to-understand. In particular, users can specify explicitly communication, synchronization, data distribution, and work distribution to facilitate performance tuning. In addition, XMP supports features for one-sided communication on each process, which was not available in HPF. This feature might enable users to implement parallel algorithms easily.

1.1 Target Hardware

The target of XcalableMP is distributed-memory multicomputers (Fig. 1). Each computation node, which may contain several cores, has its own local memory (shared by the cores, if any), and is connected with the others via an interconnection network. Each node can access its local memory directly and remote memory (the memory of another node) indirectly (i.e., through inter-node communication). However, it is assumed that accessing remote memory may be much slower than accessing local memory.

1.2 Execution Model

An XcalableMP program execution is based on the Single Program Multiple Data (SPMD) model, where each node starts execution from the same main routine, and continues to execute the same code independently (i.e., asynchronously), which is

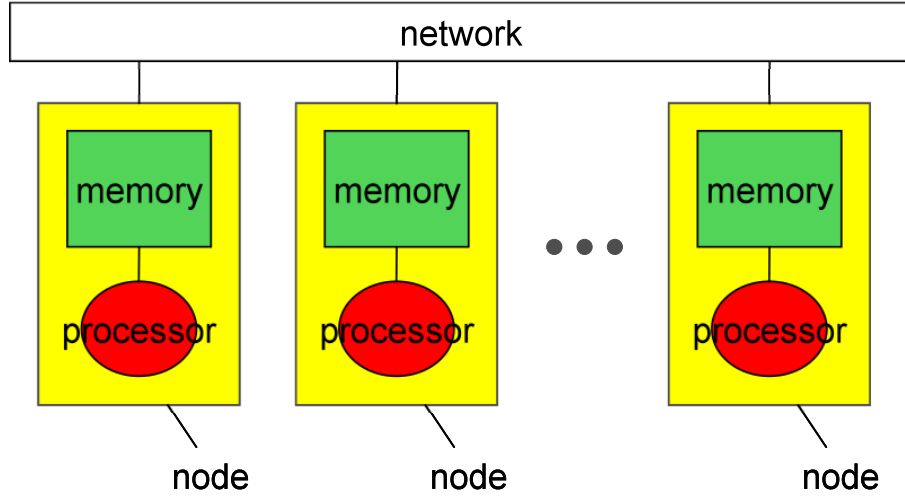


Fig. 1 Target hardware of XMP.

referred to as the *replicated execution*, until it encounters an XcalableMP construct (Fig. 2).

A set of nodes that executes a procedure, statement, loop, a block, etc. is referred to as its *executing node set*, and is determined by the innermost `task`, `loop`, or `array` directive surrounding it dynamically, or at runtime. The *current executing node set* is an executing node set of the current context, which is managed by the XcalableMP runtime system on each node.

The current executing node set at the beginning of the program execution, or *entire node set*, is a node set that contains all the available nodes, which can be specified in an implementation-defined way (e.g., through a command-line option).

When a node encounters at runtime either a `loop`, `array`, or `task` construct, and is contained by the node set specified by the `on` clause of the directive, it updates the current executing node set with the specified one and executes the body of the construct, after which it resumes the last executing node set and proceeds to execute the subsequent statements.

In particular, when a node in the current executing node set encounters a `loop` or an `array` construct, it executes the loop or the array assignment in parallel with other nodes, so that each iteration of the loop or element of the assignment is independently executed by the node in which a specified data element resides.

When a node encounters a synchronization or a communication directive, synchronization or communication occurs between it and other nodes. That is, such *global constructs* are performed collectively by the current executing nodes. Note that neither synchronization nor communication occurs unless these constructs are specified.




Fig. 2 Execution model of XMP.

1.3 Data Model

There are two classes of data in XcalableMP: *global data* and *local data*. Data declared in an XcalableMP program are local by default.

Global data are distributed onto the executing node set by the `align` directive (see section ??). Each fragment of distributed global data is allocated in the local memory of a node in the executing node set.

Local data comprises all data that are not global. They are replicated within the local memory of each of the executing nodes.

A node can access directly only local data and sections of global data that reside in its local memory. To access data in remote memory, explicit communication must be specified in such ways as global communication constructs and coarray assignments.

In particular, in XcalableMP Fortran, for common blocks that include any global variables, it is implementation-defined what storage sequences they occupy and how storage association is defined between two of them.

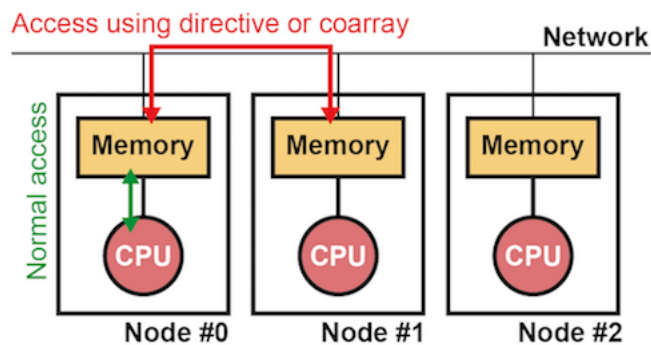


Fig. 3 Data model of XMP.

1.4 Programming Models

1.4.1 Partitioned Global Address Space

XMP can be classified as a partitioned global address space (PGAS) programming model, such as Co-Array Fortran [1], Unified Parallel C [2], and Chapel [3].

In the PGAS model, multiple executing entities (i.e. threads, processes, or nodes in XMP) share a part of their address space, which is, however, partitioned and a portion of which is local to each executing entity.

The two programming models, global-view and local-view, that XMP supports to achieve high performance and productivity on PGAS are explained below.

1.4.2 Global-view Programming Model

The global-view programming model is useful when, starting from a sequential version of a program, the programmer parallelizes it in data-parallel style by adding directives with minimum modification. In the global-view programming model, the programmer describes the distribution of data among nodes using the data distribution directives. The `loop` construct assigns each iteration of a loop to the node at which the computed data is located. The global-view communication directives are used to synchronize nodes, maintain the consistency of shadow areas, and move sections of distributed data globally. Note that the programmer must specify explicitly communications to make all data references in the program local, and this is done using appropriate directives.

In many cases, the XcalableMP program according to the global-view programming model is based on a sequential program, and it can produce the same results, regardless of the number of nodes (Fig. 4).

There are three groups of directives for the global-view programming model. Because these directives are ignored as a comment by the compilers of base languages (Fortran and C), an XcalableMP program can be compiled by them to ensure that they run properly.

- *Data mapping*, which specifies the data distribution and mapping to nodes (partially inherited from HPF).
- *Work mapping (parallelization)*, which assigns a work to a node set. The `loop` construct maps each iteration of a loop to nodes owning a specific data elements. The `task` construct defines a set amount of work as a *task*, and assigns it to a specific node set.
- *Communication and synchronization*, which specify how to communicate and synchronize with the other compute nodes. In XcalableMP, inter-node communication must be explicitly specified by the programmer. The compiler guarantees that no communication occurs unless it is explicitly specified by the programmer.

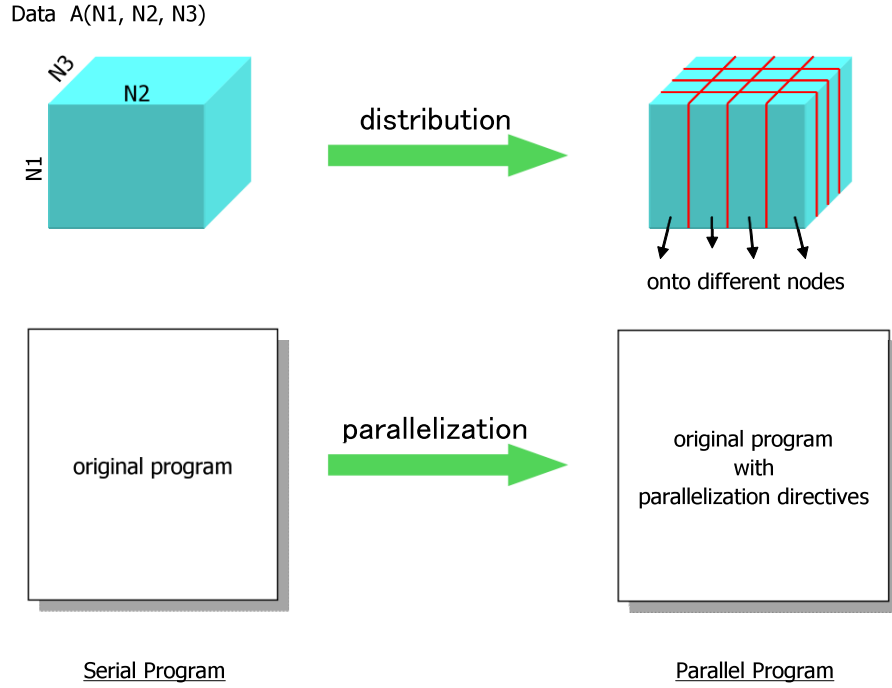


Fig. 4 Parallelization using the global-view programming model.

1.4.3 Local-view Programming Model

The local-view programming model is suitable for programs that implement an algorithm and a remote data reference that are to be executed by each node (Fig. 5).

For the local-view programming model, some language extensions and directives are provided. The coarray notation, which is imported from Fortran 2008, is one such extension, and can be used to specify data on which node

is to be accessed. For example, the expression of $A(i)[N]$ is used to access an array element of $A(i)$ located on the node N . If the access is a reference, then a one-sided communication to read the value from the remote memory (i.e., the *get* operation) is issued by the executing node. If the access is a definition, then a one-sided communication to write the value to the remote memory (i.e., the *put* operation) is issued by the executing node.

1.4.4 Mixture of Global View and Local View

In the global-view model, nodes are used to distribute data and works. In the local-view model, nodes are used to address remote data in the coarray notation. In application programs, programmers should choose an appropriate data model ac-

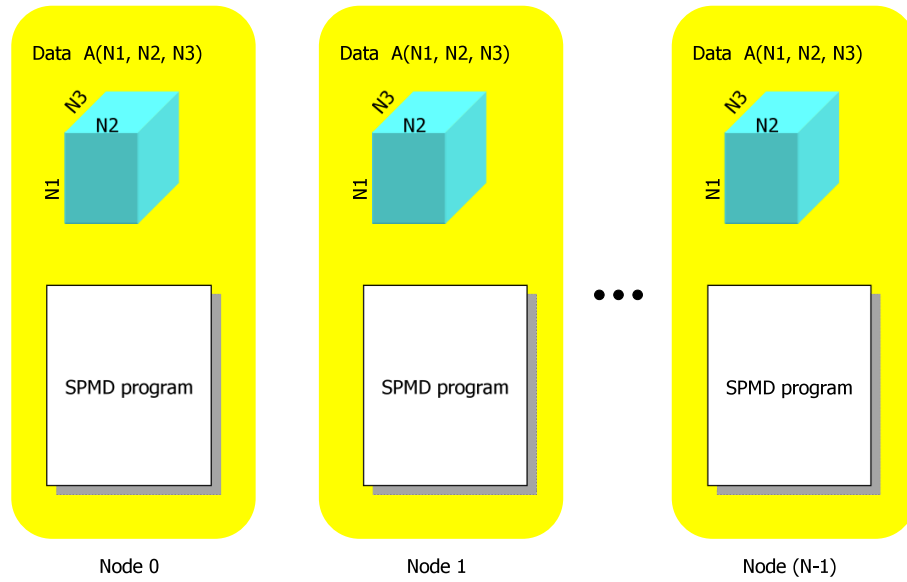


Fig. 5 Local-view programming model.

cording to the characteristics of their program. Fig. 6 illustrates the global view and the local view of data.

Data may have both a global view and a local view, and can be accessed in both of the views. XcalableMP provides a directive to give the local name (alias) to global data declared in the global-view programming model to enable them to also be accessed in the local-view programming model. This feature is useful to optimize a certain part of a program by using explicit remote data access in the local-view programming model.

1.5 Base Languages

The XcalableMP language specification is defined based on Fortran and C as the base languages. More specifically, the base language of XcalableMP Fortran is Fortran 90 or later, and that of XcalableMP C is ISO C90 (ANSI C89) or later with some extensions (see below).

1.5.1 Array Section in XcalableMP C

In XcalableMP C, the base language C is extended so that a part of an array, that is, an *array section* or *subarray*, can be put in an *array assignment statement*, which

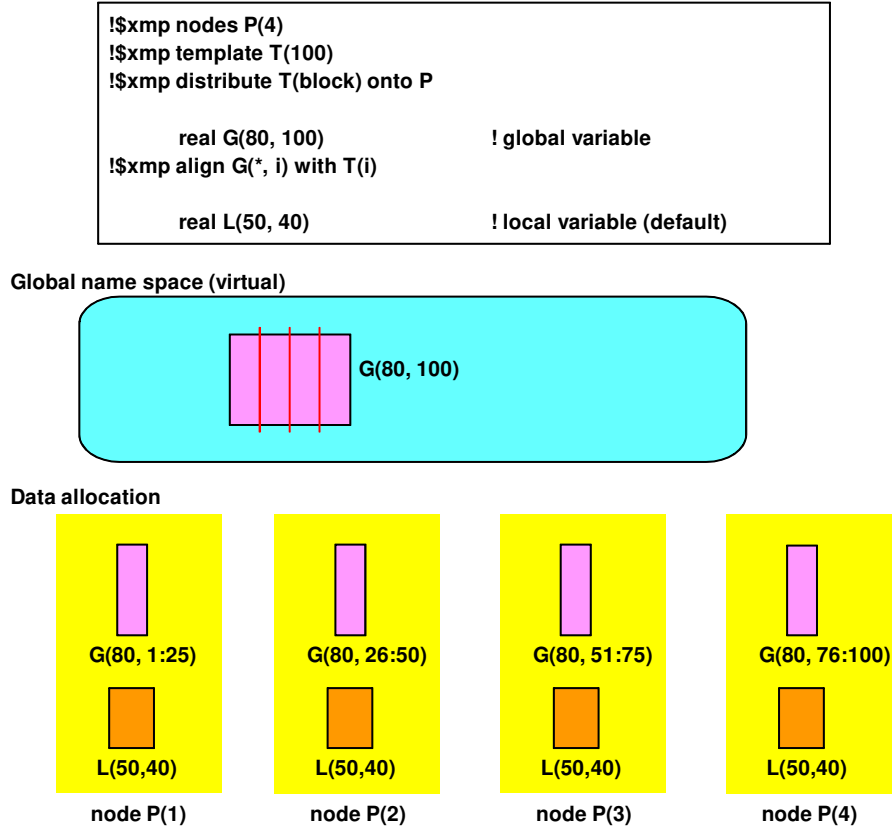


Fig. 6 Global view and local view.

is described in 1.5.2, and some XcalableMP constructs. An array section is built from a subset of the elements of an array, which is specified by a sequence of square-bracketed integer expressions or *triplets*, which are in the form of:

$$[\textit{base}] : [\textit{length}] [: \textit{step}]$$

When *step* is positive, the *triplet* specifies a set of subscripts that is a regularly spaced integer sequence of length *length* beginning with *base* and proceeding in increments of *step* up to the largest. The same applies to negative *step* too.

When *base* is omitted, it is assumed to be 0. When *length* is omitted, it is assumed to be the number of remainder elements of the dimension of the array. When *step* is omitted, it is assumed to be 1.

Example

Assuming that an array *A* is declared by the following statement,


```
int A[100];
```

some array sections can be specified as follows:

```
A[10:10]  array section of 10 elements from A[10] to A[19]
A[10:]    array section of 90 elements from A[10] to A[99]
A[:10]    array section of 10 elements from A[0] to A[9]
A[10:5:2] array section of 5 elements from A[10] to A[18] by step 2
A[:]      the whole of A
```

1.5.2 Array Assignment Statement in XscalableMP C

The value of each element of the result of the right-hand side expression is assigned to the corresponding element of the array section on the left-hand side. When an operator or an elemental function (see section 1.5.3.2) is applied to array sections in the right-hand side expression, it is evaluated to an array section that has the same shape as that of the operands or arguments, and each element of which is the result of the operator or function applied to the corresponding element of the operands or arguments. A scalar object is assumed to be an array section that has the same shape as that of the array section(s), and where each element has its value.

Note that an array assignment is a statement, and therefore cannot appear as an expression in any other statements.

Examples

An array assignment statement in the fourth line copies the elements B[0] through B[4] into the elements A[5] through A[9].

XscalableMP C
<pre>int A[10]; int B[5]; ... A[5:5] = B[0:5];</pre>

1.5.3 Support for Intrinsic/Built-in Functions of the Base Languages

This section describes how intrinsic/built-in functions of the base languages work on XMP's global arrays.

Many other intrinsic and library procedures, such as mapping inquiry functions and transformational procedures, are defined in XMP. See the language specification for their detail.

1.5.3.1 Array Intrinsic Functions in XcalableMP Fortran

The array intrinsic functions of the base language Fortran are classified into three classes: *inquiry*, *elemental*, and *transformational*.

It is specified as follows how these functions work in the XMP/F programs when a global array appears as an argument.

- Inquiry functions
The inquiry functions with a global array or its subobject being an argument are regarded as inquiries about the global array, and return its “global” properties as if it were not distributed.
- Elemental functions
The result of the elemental functions with a global array or its subobject being an argument has the same shape and mapping as the argument. Note that such a reference of these elemental functions is in effect limited to be in the array construct.
- Transformational functions
It is unspecified how the transformational functions work when a global array or its subobject appears as an argument. A processor shall detect such a reference of these functions and issue a warning message for it. Some intrinsic transformational subroutines are defined in section ?? as alternatives to these transformational functions.

1.5.3.2 Built-in Elemental Functions in XcalableMP C

Some built-in elemental functions that can operate each element of array arguments are defined in XcalableMP C. Such a built-in function accepts one or more array sections as its arguments and returns an array-valued result having the same shape and mapping as the argument. The values of the elements of the result are the same as what would have been obtained if the scalar function of the C standard library had been applied separately to the corresponding elements of each array argument.

These functions may appear on the right-hand side of an array assignment statement, and it should be preceded by the `array` directive if the array section is distributed.

Table 1 shows the list of built-in elemental functions in XcalableMP C. Their elementwise behavior is the same as those of the corresponding functions in the C standard library.

1.6 Interoperability

Most of existing parallel applications are written with MPI. It is not realistic to port them over to XMP because each of them consists of millions of lines.

Table 1 Built-in elemental functions in XcalableMP C. (The first line refers to the element type of their argument(s) and return value.)

double	float	long double
acos	acosf	acosl
asin	asinf	asinl
atan	atanf	atanl
atan2	atan2f	atan2l
cos	cosf	cosl
sin	sinf	sinl
tan	tanf	tanl
cosh	coshf	coshl
sinh	sinhf	sinhl
tanh	tanhf	tanhl
exp	expf	expl
frexp	frexpf	frexpl
ldexp	ldexpf	ldexpl
log	logf	logl
log10	log10f	log10l
fabs	fabsf	fabsl
pow	powf	powl
sqrt	sqrtf	sqrtl
ceil	ceilf	ceill
floor	floorf	floorl
fmod	fmodf	fmodl

Because XMP is interoperable with MPI, users can develop an XMP application by modifying a part of an existing one instead of rewriting it totally. Besides, when developing a parallel application from scratch, it is possible to use XMP to write a complicated part of, for example, domain decomposition while they use MPI, which could be faster than XMP, to write a hot-spot part that need to be tuned carefully. In addition, XMP is interoperable with OpenMP and Python (see Chapter ??).

It might be difficult to develop an application with just one programming language or framework since it generally has its own strong and weak points. Thus, an XMP program is interoperable with those in other languages to provide both high productivity and performance.

2 Data Mapping

2.1 nodes Directive

The nodes directive declares the name of a node array and its shape. A node array can be multi-dimensional.

XcalableMP C
#pragma xmp nodes p[4]
XcalableMP Fortran
!\$xmp nodes p(4)

The nodes directive declares 1-dimensional node set `p` which has four nodes. In XMP/C, the node set consists of `p[0]`, `p[1]`, `p[2]`, and `p[3]`. In XMP/Fortran, the node set consists of `p(1)`, `p(2)`, `p(3)`, and `p(4)`.

2.1.1 Multi-dimensional Node Set

XcalableMP C
#pragma xmp nodes p[2][3]
XcalableMP Fortran
!\$xmp nodes p(3,2)

The nodes directive declares two-dimensional node array `p` that includes six nodes. In XMP/C, the node set includes `p[0][0]`, `p[0][1]`, `p[0][2]`, `p[1][0]`, `p[1][1]`, and `p[1][2]`. In XMP/Fortran, the node set includes `p(1,1)`, `p(2,1)`, `p(3,1)`, `p(1,2)`, `p(2,2)`, and `p(3,2)`.

Note: The ordering of the elements in a node array depends on the base language, that is, C or Fortran.

XcalableMP C
#pragma xmp nodes p[*]
XcalableMP Fortran
!\$xmp nodes p(*)

An asterisk can be specified in the nodes directive to declare a *dynamic* node array. In the above program, one-dimensional dynamic node array `p` is declared with an asterisk. The size of a dynamic node array is determined at runtime (at the beginning of the execution). For example, when the user runs the sample program with three nodes, the node array `p` include three nodes.

The user also can declare multi-dimensional dynamic nodes with an asterisk.

XcalableMP C
#pragma xmp nodes p[*][3]
XcalableMP Fortran
!\$xmp nodes p(3,*)

When the user runs the sample program with 12 nodes, the node array `p` will have a shape of 4x3, in C, or 3x4, in Fortran.

Note: The user can put an asterisk only in the last dimension of the node array.

Hint: The dynamic node array may interfere with compiler optimizations. In general, static one achieve better performance.

The user can declare a node subarray derived from the existing node array. Node subarrays can be used, for example, to optimize inter-node communication by reducing the number of nodes participating in the communication.

XcalableMP C

```
#pragma xmp nodes p[16]
#pragma xmp nodes q[8]=p[0:8]
#pragma xmp nodes r[4][2]=p[8:8]
```

XcalableMP Fortran

```
!$xmp nodes p(16)
!$xmp nodes q(8)=p(1:8)
!$xmp nodes r(2,4)=p(9:16)
```

In line 1, a node array `p` including 16 nodes is declared. In line 2, a node subarray `q` derived from the first half of `p` is declared. In line 3, a two-dimensional node subarray `r` derived from the latter half of `p` is declared.

The user can declare a n-dimensional node subarray derived from a m-dimensional one.

XcalableMP C

```
#pragma xmp nodes p[4][2]
#pragma xmp nodes row[4]=p[:,*]
#pragma xmp nodes col[2]=p[*][:]
```

XcalableMP Fortran

```
!$xmp nodes p(2,4)
!$xmp nodes row(4)=p(*,:)
!$xmp nodes col(2)=p(:,*)
```

In line 1, a two-dimensional node array `p` including 4x2 nodes is declared. In line 2, a node subarray `row` derived from a single row of `p` is declared. In line 3, a node subarray `col` derived from a single column of `p` is declared.

A colon represents a triplet which indicate all possible indices in the dimension. An asterisk indicate the index of the current executing node in the dimension. For example, `col[2]` is `p[0][0:2]` on nodes `p[0][0]` and `p[0][1]`, and is `p[1][0:2]` on nodes `p[1][0]` and `p[1][1]` in XMP/C. Similarly, `col(2)` is `p(1:2, 1)` on nodes `p(1, 1)` and `p(2, 1)`, and is `p(1:2, 2)` on nodes `p(1, 2)` `p(2, 2)` in XMP/Fortran.

In XMP/C, both `p[0][0]` and `p[0][1]` will be `row[0]`. Likewise, `p[0][0]`, `p[1][0]`, `p[2][0]` and `p[3][0]` will be `col[0]` in each execution context. In XMP/Fortran, both

XMP/C

p[0][0]	p[0][1]
p[1][0]	p[1][1]
p[2][0]	p[2][1]
p[3][0]	p[3][1]
row[4] row[4]	

p[0][0]	p[0][1]	col[2]
p[1][0]	p[1][1]	col[2]
p[2][0]	p[2][1]	col[2]
p[3][0]	p[3][1]	col[2]

XMP/Fortran

p(1,1)	p(2,1)
p(1,2)	p(2,2)
p(1,3)	p(2,3)
p(1,4)	p(2,4)
row[4] row[4]	

p(1,1)	p(2,1)	col[2]
p(1,2)	p(2,2)	col[2]
p(1,3)	p(2,3)	col[2]
p(1,4)	p(2,4)	col[2]

Fig. 7 Partial node set.

p(1,1) and p(2,1) will be row(1). Likewise, p(1,1), p(1,2), p(1,3) and p(1,4) will be col(1) in each context.

Note: The semantics of an asterisk in a node reference is different from that in a declaraiion.

2.2 template Directive

The `template` directive declares the name and shape of a template.

_____ XcalableMP C _____
#pragma xmp template t[10]
_____ XcalableMP Fortran _____
!\$xmp template t(10)

The `template` directive declares a one-dimensional template `t` having ten elements. Templates are indexed in the similar manner to arrays in the base languages. For the above examples, the template `t` is indexed from zero to nine (i.e. `t[0] ... t[9]`) in XMP/C while from one to ten (i.e. `t(1) ... t(10)`) in XMP/Fortran.

Hint: In many cases, a template should be declared to have the same shape as your target array.

	XcalableMP C	
#pragma xmp template t[10][20]		
	XcalableMP Fortran	
!\$xmp template t(20,10)		

The `template` directive declares a two-dimensional template `t` that has 10x20 elements. In XMP/C, `t` is indexed from `t[0][0]` to `t[9][19]` while, in XMP/Fortran, from `t(1,1)` to `t(20,10)`.

	XcalableMP C	
#pragma xmp template t[:]		
	XcalableMP Fortran	
!\$xmp template t(:)		

In the above examples, a colon instead of an integer is specified as the size to declare a one-dimensional dynamic template `t`. The colon indicates that the size of the template is not fixed and to be fixed at runtime by the `template_fix` construct (Sec. 2.6).

2.3 distribute Directive

The `distribute` directive specifies a distribution of the target template. The user can specify a distribution format of block, cyclic, block-cyclic, or gblock (i.e. uneven block).

2.3.1 Block Distribution

	XcalableMP C	
#pragma xmp distribute t[block] onto p		
	XcalableMP Fortran	
!\$xmp distribute t(block) onto p		

The target template is divided into contiguous blocks and distributed among nodes. Let's suppose that the size of the template is N and the number of nodes is


K . If N is divisible by K , a block of size N/K are assigned to each node; otherwise a block of size $\text{ceil}(N/K)$ is assigned to each of $N/\text{ceil}(N/K)$ nodes, a block of size $\text{mod}(N, K)$ to one node, and no block to $(K - N/\text{ceil}(N/K) - 1)$ nodes. The block distribution is useful for stencil computation.

Note: The function $\text{ceil}(x)$ returns a minimum integer value greater than x , and $\text{mod}(x, y)$ returns x modulo y .



XcalableMP C
<pre>#pragma xmp nodes p[3] #pragma xmp template t[22] #pragma xmp distribute t[block] onto p</pre>

XcalableMP Fortran
<pre>!\$xmp nodes p(3) !\$xmp template t(22) !\$xmp distribute t(block) onto p</pre>

XMP/C



XMP/Fortran



Since $\text{ceil}(22/3)$ is 8, eight elements are allocated on each of $p[0]$ and $p[1]$. And then, remaining six elements are allocated on $p[2]$.

2.3.2 Cyclic Distribution

XcalableMP C
<pre>#pragma xmp distribute t[cyclic] onto p</pre>

XcalableMP Fortran
<pre>!\$xmp distribute t(cyclic) onto p</pre>


The target template is divided into chunks of size one and distributed among nodes in a round-robin manner. The cyclic distribution is usefull for the case where the load on each element of the template is not balanced.

XcalableMP C

```
#pragma xmp nodes p[3]
#pragma xmp template t[22]
#pragma xmp distribute t[cyclic] onto p
```

XcalableMP Fortran

```
!$xmp nodes p(3)
!$xmp template t(22)
!$xmp distribute t(cyclic) onto p
```

XMP/C**XMP/Fortran****2.3.3 Block-cyclic Distribution**

XcalableMP C

```
#pragma xmp distribute t[cyclic(w)] onto p
```

XcalableMP Fortran

```
!$xmp distribute t(cyclic(w)) onto p
```




The target template is divided into chunks of size w and distributed among nodes in a round-robin manner. The block-cyclic distribution is useful for the case where the load on each element of the template is not balanced but the locality of the elements is required.

XcalableMP C

```
#pragma xmp nodes p[3]
#pragma xmp template t[22]
#pragma xmp distribute t[cyclic(3)] onto p
```

XcalableMP Fortran

```
!$xmp nodes p(3)
!$xmp template t(22)
!$xmp distribute t(cyclic(3)) onto p
```

XMP/C**XMP/Fortran****2.3.4 Gblock Distribution**

```

XcalableMP C
#pragma xmp distribute t[gblock(W)] onto p

```

```

XcalableMP Fortran
!$xmp distribute t(gblock(W)) onto p

```

The target template is divided into contiguous blocks of size $W[0]$, $W[1]$, \dots , in XMP/C, or $W(1)$, $W(2)$, \dots , in XMP/Fortran, and distributed among nodes. An array W is called a mapping array. The user can specify irregular (uneven) block distribution with `gblock`.

```



XcalableMP C
#pragma xmp nodes p[3]
#pragma xmp template t[22]
int W[3] = {6, 11, 5};
#pragma xmp distribute t[gblock(W)] onto p

```

```

XcalableMP Fortran
!$xmp nodes p(3)
!$xmp template t(22)
integer, parameter :: W(3) = (/6,11,5/)
!$xmp distribute t(gblock(W)) onto p

```

XMP/C**XMP/Fortran**

The user can specify an asterisk instead of a mapping array in the `gblock` distribution to defer fixing the actual distribution. In such a case, the actual distribution will be fixed at runtime by using `template_fix` construct.

2.3.5 Distribution of Multi-dimensional Templates

The user can distribute a multi-dimensional template onto a node array.

```

XcalableMP C
#pragma xmp nodes p[2][2]
#pragma xmp template t[10][10]
#pragma xmp distribute t[block][block] onto p

```


```

XcalableMP Fortran
!$xmp nodes p(2,2)
!$xmp template t(10,10)
!$xmp distribute t(block,block) onto p


```

The distribute directive declares the distribution of a two-dimensional template `t` on a two-dimensional node array `p`. Each dimension of the template is divided in block and each rectangular region is assigned to a node.

XMP/C



XMP/Fortran



The user can specify a different distribution format in each of the dimension of a template.

```

XcalableMP C
#pragma xmp nodes p[2][2]
#pragma xmp template t[10][10]
#pragma xmp distribute t[block][cyclic] onto p



```

```

XcalableMP Fortran
!$xmp nodes p(2,2)
!$xmp template t(10,10)
!$xmp distribute t(cyclic,block) onto p

```




When an asterisk is specified in the distribute directive as a distribution format, the target dimension is “non-distributed.” In the following example, the first dimension will be distributed in a block manner and the second dimension will be non-distributed.

XMP/C**XMP/Fortran****XcalableMP C**

```
#pragma xmp nodes p[4]
#pragma xmp template t[10][10]
#pragma xmp distribute t[block][*] onto p
```

XcalableMP Fortran

```
!$xmp nodes p(4)
!$xmp template t(10,10)
!$xmp distribute t(*,block) onto p
```

XMP/C**XMP/Fortran**

2.4 align Directive

The align directive specifies that an array is to be mapped in the same way as a specified template.

XcalableMP C

```
#pragma xmp nodes p[4]
#pragma xmp template t[8]
#pragma xmp distribute t[block] onto p
```

```

int a[8];
5 #pragma xmp align a[i] with t[i]

```


XscalableMP Fortran

```

!$xmp nodes p(4)
!$xmp template t(8)
!$xmp distribute t(block) onto p
integer :: a(8)
5 !$xmp align a(i) with t(i)

```

The array a is decomposed and laid out so that each array element $a(i)$ is colocated with the corresponding template element $t(i)$.



The align directive can also be used for multi-dimensional arrays.

XscalableMP C

```

#pragma xmp nodes p[2][2]
#pragma xmp template t[8][8]
#pragma xmp distribute t[block][block] onto p
int a[8][8];
5 #pragma xmp align a[i][j] with t[i][j]

```

XscalableMP Fortran

```

!$xmp nodes p(2,2)
!$xmp template t(8,8)
!$xmp distribute t(block,block) onto p
integer :: a(8,8)
5 !$xmp align a(j,i) with t(j,i)

```


The user can align a two-dimensional array with a one-dimensional template.

XscalableMP C

```

#pragma xmp nodes p[4]
#pragma xmp template t[8]
#pragma xmp distribute t[block] onto p
int a[8][8];
5 #pragma xmp align a[i][*] with t[i]

```




```

XcalableMP Fortran
!$xmp nodes p(4)
!$xmp template t(8)
!$xmp distribute t(block) onto p
integer :: a(8,8)
5 !$xmp align a(*,i) with t(i)

```

When an asterisk is specified as a subscript in a dimension of the target array in the `align` directive, the dimension is not distributed. In the sample program above, the first dimension of the array `a` is distributed onto the node array `p` while the second dimension is “collapsed.”



In XMP/C, `a[0:2][:]` will be allocated on `p[0]` while, in XMP/Fortran, `a(:, 1:2)` will be allocated on `p(1)`.

The user also can align a one-dimensional array with a two-dimensional template.

```

XcalableMP C
#pragma xmp nodes p[2][2]
#pragma xmp template t[8][8]
#pragma xmp distribute t[block][block] onto p
int a[8];
5 #pragma xmp align a[i] with t[i][*]

XcalableMP Fortran
!$xmp nodes p(2,2)
!$xmp template t(8,8)


```

```

!$xmp distribute t(block,block) onto p
integer :: a(8)
5 !$xmp align a(i) with t(*,i)

```

When an asterisk is specified as a subscript in a dimension of the target template in the align directive, the array will be “replicated” along the axis of the dimension.



In XMP/C, $a[0:4]$ will be replicated and allocated on $p[0][0]$ and $p[0][1]$ while, in XMP/Fortran, $a(1:4)$ will be allocated on $p(1,1)$ and $p(2,1)$.

2.5 Dynamic Allocation of Distributed Array

This section explains how distributed (i.e. global) arrays are allocated at runtime. The basic procedure is common in XMP/C and XMP/Fortran with a few specific difference.

```

XcalableMP C
#pragma xmp nodes p[4]
#pragma xmp template t[N]
#pragma xmp distribute t[block] onto p
float *a;
5 #pragma xmp align a[i] with t[i]
:
a = xmp_malloc(xmp_desc_of(a), N);

```

In XMP/C, first, declare a pointer of the type of the target array; second, align it as if it were an array; finally, allocate memory for it with the `xmp_malloc()` function. `xmp_desc_of()` is an intrinsic/builtin function that returns the descriptor of the XMP object specified by the argument.

```

XcalableMP Fortran
!$xmp nodes p(4)
!$xmp template t(N)
!$xmp distribute t(block) onto p

```

```

real, allocatable :: a(:)
5 !$xmp align a(i) with t(i)

allocate(a(N))

```

In XMP/Fortran, first, declare an allocatable array; Second, align it; finally, allocate memory for it with the allocate statement.

For multi-dimensional arrays, the procedure is the same as that for one-dimensional arrays.

```

XcalableMP C
#pragma xmp nodes p[2][2]
#pragma xmp template t[N1][N2]
#pragma xmp distribute t[block][block] onto p
float (*a)[N2];
5 #pragma xmp align a[i][j] with t[i][j]
:
a = (float (*)[N2])xmp_malloc(xmp_desc_of(a), N1, N2);

```

```

XcalableMP Fortran
!$xmp nodes p(2,2)
!$xmp template t(N2,N1)
!$xmp distribute t(block,block) onto p
real, allocatable :: a(:, :)
5 !$xmp align a(j,i) with t(j,i)
:
allocate(a(N2,N1))

```

Note: If the size of the template is not fixed until runtime, you have to use the `template_fix` construct.

2.6 template_fix Construct

The `template_fix` construct fixes the shape and/or the distribution of an undefined template.

```

XcalableMP C
#pragma xmp nodes p[4]
#pragma xmp template t[:]
#pragma xmp distribute t[block] onto p
double *a;
5 #pragma xmp align a[i] with t[i]

```



```

int n = 100;
#pragma xmp template_fix t[n]
a = xmp_malloc(xmp_desc_of(a), n);

```

XcalableMP Fortran

```

!$xmp nodes p(4)
!$xmp template t(:)
!$xmp distribute t(block) onto p
real, allocatable :: a(:)
5 integer :: n
!$xmp align a(i) with t(i)

n = 100
!$xmp template_fix t(n)
10 allocate(a(n))

```

In the above sample code, first, a template `t` whose size is undefined (“:”) is declared; second, a pointer `a` in XMP/C or an allocatable array `a` in XMP/Fortran is aligned with the template; third, the size of the template is fixed with a `template_fix` construct; finally, the pointer or the allocatable array is allocated with the `xmp_malloc()` builtin function in XMP/C or the `allocate` statement in XMP/Fortran, respectively.

Note: `template_fix` constructs can be applied to a template only once.

This construct can also be used to fix a mapping array of a template that is distributed in “`gblock(*)`” at declaration.

XcalableMP C

```

#pragma xmp nodes p[4]
#pragma xmp template t[:]
#pragma xmp distribute t[gblock(*)] onto p
double *a;
5 #pragma xmp align a[i] with t[i]

int n = 100;
int m[] = {40,30,20,10};

10 #pragma xmp template_fix[gblock(m)] t[n]
a = xmp_malloc(xmp_desc_of(a), n);

```

XcalableMP Fortran

```

!$xmp nodes p(4)
!$xmp template t(:)
!$xmp distribute t(gblock) onto p

```

```

real, allocatable :: a(:)
5 integer :: n, m(4)
!$xmp align a(i) with t(i)

n = 100
m(:) = (/40,30,20,10/)
10 !$xmp template_fix(gblock(m)) t(n)
allocate(a(n))

```

3 Work Mapping

3.1 task and tasks Construct

The task construct generates a task, which is executed by the specified nodes. The tasks construct asserts that surrounding task constructs can be executed in parallel.

3.1.1 task Construct

The on clause of the task construct specifies the node set that executes the task.

```

_____ XcalableMP C _____
#include <stdio.h>
#pragma xmp nodes p[4]

int main(){
5   int num = xmpc_node_num();
#pragma xmp task on p[1:3]
{
    printf("%d: Hello\n", num);
}
10
    return 0;
}

```

```

_____ XcalableMP Fortran _____
program main
!$xmp nodes p(4)
    integer :: num

5   num = xmp_node_num()
!$xmp task on p(2:4)
    write(*,*) num, ": Hello"
!$xmp end task

```

```
10 end program main
```

In the above example, nodes `p[1]`, `p[2]`, and `p[3]` invokes the `printf()` function, and `p[1]` outputs “1: Hello” in XMP/C; `p(2)`, `p(3)`, and `p(4)` execute the `write` statement, and `p(2)` outputs “2: Hello” in XMP/Fortran.

Note that a new node set is generated by each task construct. Let’s consider inserting a `bcast` construct into the task.

XcalableMP C

```
#pragma xmp task on p[1:3]
{
  #pragma xmp bcast (num)
}
```

XcalableMP Fortran

```
!$xmp task on p(2:4)
!$xmp bcast (num)
!$xmp end task
```

This `bcast` construct is executed by the node set specified by the task construct. Thus, the node `p[1]` broadcasts the value to `p[2]` and `p[3]` in XMP/C, and `p(2)` to `p(3)` and `p(4)` in XMP/Fortran.

XMP/C

	p[0]	p[1]	p[2]	p[3]	
num =	0	1	2	3	
num =	0	1	1	1	} Region by task directive
		bcast directive			
num =	0	1	1	1	

XMP/Fortran

	p(1)	p(2)	p(3)	p(4)	
num =	1	2	3	4	
num =	1	2	2	2	} Region by task directive
		bcast directive			
num =	1	2	2	2	

The `bcast` construct in the above code is equivalent to that in the following code, where it is executed by a new node set that is explicitly declared.

<pre> XcalableMP C #pragma xmp nodes q[3] = p[1:3] #pragma xmp bcast (num) on q </pre>
<pre> XcalableMP Fortran !\$xmp nodes q(3) = p(2:4) !\$xmp bcast (num) on q </pre>

Note that the task is executed by the node set specified by the `on` clause. Therefore, `xmpc_node_num()` and `xmp_node_num()` return the id in the node set.

For example, consider inserting `xmpc_node_num()` or `xmp_node_num()` into the task in the first program.

<pre> XcalableMP C #include <stdio.h> #pragma xmp nodes p[4] int main(){ 5 #pragma xmp task on p[1:3] { printf("%d: Hello\n", xmpc_node_num()); } 10 return 0; } </pre>
<pre> XcalableMP Fortran program main !\$xmp nodes p(4) !\$xmp task on p(2:4) 5 write(*,*) xmp_node_num(), ": Hello" !\$xmp end task end program main </pre>

The node `p[1]` outputs “0: Hello” in XMP/C, and `p(2)` “1: Hello” in XMP/Fortran.

Note: A new node set should be collectively generated by all of the executing nodes at the point of a `task` construct unless it is surrounded by a `tasks` construct. In the above example, `p[0]` in XMP/C and `p(1)` in XMP/Fortran must execute the task construct.

3.1.2 tasks Construct

Let's consider that each of two tasks invokes a function.

	XcalableMP C
<pre> #pragma xmp nodes p[4] #pragma xmp task on p[0:2] { 5 func_a(); } #pragma xmp task on p[2:2] { 10 func_b(); } </pre>	<pre> #pragma xmp nodes p[4] #pragma xmp task on p[0:2] { func_a(); } #pragma xmp task on p[2:2] { func_b(); } </pre>



	XcalableMP Fortran
<pre> !\$xmp nodes p(4) !\$xmp task on p(1:2) call func_a() 5 !\$xmp end task !\$xmp task on p(3:4) call func_b() 10 !\$xmp end task </pre>	<pre> !\$xmp nodes p(4) !\$xmp task on p(1:2) call func_a() !\$xmp end task !\$xmp task on p(3:4) call func_b() !\$xmp end task </pre>

In the above example, the two tasks cannot be executed in parallel because those on clauses must be evaluated by all of the executing nodes.

Use the tasks construct to execute multiple tasks in parallel.

	XcalableMP C
<pre> #pragma xmp nodes p[4] #pragma xmp tasks { 5 #pragma xmp task on p[0:2] { func_a(); } #pragma xmp task on p[2:2] 10 { func_b(); } } </pre>	<pre> #pragma xmp nodes p[4] #pragma xmp tasks { #pragma xmp task on p[0:2] func_a(); #pragma xmp task on p[2:2] func_b(); } </pre>

	XcalableMP Fortran
<pre> !\$xmp nodes p(4) !\$xmp tasks </pre>	<pre> !\$xmp nodes p(4) !\$xmp tasks </pre>

XMP/C**XMP/Fortran**

```

!$xmp task on p(1:2)
5  call func_a()
!$xmp end task
!$xmp task on p(3:4)
  call func_b()
!$xmp end task
10 !$xmp end tasks



```

Because the node sets specified by the `on` clauses of the `task` constructs surrounded by a `tasks` construct are disjoint, they can be executed in parallel.

3.2 loop Construct

The `loop` construct is used to parallelize a loop. Distributed arrays in such a loop must fulfill the following two conditions:

1. There is no data/control dependence among the iterations. In other words, the iterations of the loop can be executed in any order to produce the same result.

XMP/C**XMP/Fortran**

2. An element of a distributed array is accessed only by the node that owns the element.

3.2.1 Accessing Distributed Array

The programs below are examples of a right loop directive and a loop statement. The condition 1. is satisfied because `i` is the only one index of the distributed array `a` that is accessed within the loop, and the condition 2 is also satisfied because the indices of the template in the `on` clause of the loop directive is identical to that of the distributed array.

```

XcalableMP C
#pragma xmp nodes p[2]
#pragma xmp template t[10]
#pragma xmp distribute t[block] onto p

5 int main(){
    int a[10];
    #pragma xmp align a[i] with t[i]

    #pragma xmp loop on t[i]

```

```

10  for(int i=0;i<10;i++)
    a[i] = i;

    return 0;
}

```

XcalableMP Fortran

```


program main
!$xmp nodes p(2)
!$xmp template t(10)
!$xmp distribute t(block) onto p
5  integer a(10)
!$xmp align a(i) with t(i)

!$xmp loop on t(i)
do i=1, 10
10  a(i) = i
enddo


end program main

```

XMP/C



XMP/Fortran



Is it possible to parallelize the below loops whose bounds are shrunk?

XcalableMP C

```

#pragma xmp loop on t[i]
for(int i=1;i<9;i++)
    a[i] = i;

```

XcalableMP Fortran



```

!$xmp loop on t(i)
do i=2, 9
    a(i) = i
enddo

```

In this case, the conditions 1 and 2 are satisfied and therefore it is possible to parallelize them. In XMP/C, $p[0]$ processes the indices from one to four and $p[1]$

from five to eight. In XMP/Fortran, $p(1)$ processes the indices from two to five and $p(2)$ from six to nine.

XMP/C**XMP/Fortran**

Next, is it possible to parallelize the below loops in which the index of the distributed array is different?



XcalableMP C

```
#pragma xmp loop on t[i]
for(int i=1;i<9;i++)
    a[i+1] = i;
```

XcalableMP Fortran

```
!$xmp loop on t(i)
do i=2, 9
    a(i+1) = i
enddo
```

In this case, the condition 1 is satisfied but 2 is not, and therefore it is not possible to parallelize them. In XMP/C, $p[0]$ tries to access $a[5]$ but does not own it. In XMP/Fortran, $p(1)$ tries to access $a(6)$ but does not own it.

XMP/C**XMP/Fortran**

3.2.2 Reduction Computation

The serial programs below are examples of the reduction computation.

C

```

#include <stdio.h>

int main(){
    int a[10], sum = 0;

5   for(int i=0;i<10;i++){
        a[i] = i+1;
        sum += a[i];
    }

10  printf("%d\n", sum);

    return 0;
}

```

Fortran

```

program main
    integer :: a(10), sum = 0

5   do i=1, 10
        a(i) = i
        sum = sum + a(i)
    enddo

10  write(*,*) sum

end program main

```

If the above loops are parallelized only with the loop directive, the value of the variable `sum` varies from node to node because it is calculated separately on each node.

XcalableMP C

```

5   #pragma xmp loop on t[i]
    for(int i=0;i<10;i++){
        a[i] = i+1;
        sum += a[i];
    }

```

XcalableMP Fortran

```



!$xmp loop on t(i)
5   do i=1, 10
        a(i) = i
    enddo

```

```

    sum = sum + a(i)
5  enddo

```

XMP/C**XMP/Fortran**

Then, add the reduction clause to the loop directive.

XcalableMP C

```

#include <stdio.h>
#pragma xmp nodes p[2]
#pragma xmp template t[10]
#pragma xmp distribute t[block] onto p
5
int main(){
    int a[10], sum = 0;
    #pragma xmp align a[i] with t[i]
10 #pragma xmp loop on t[i] reduction(+:sum)
    for(int i=0;i<10;i++){
        a[i] = i+1;
        sum += a[i];
    }
15
    printf("%d\n", sum);

    return 0;
}

```

XcalableMP Fortran

```

program main
!$xmp nodes p(2)
!$xmp template t(10)
!$xmp distribute t(block) onto p
5   integer :: a(10), sum = 0
!$xmp align a(i) with t(i)

!$xmp loop on t(i) reduction(+:sum)
   do i=1, 10
10      a(i) = i
        sum = sum + a(i)

```

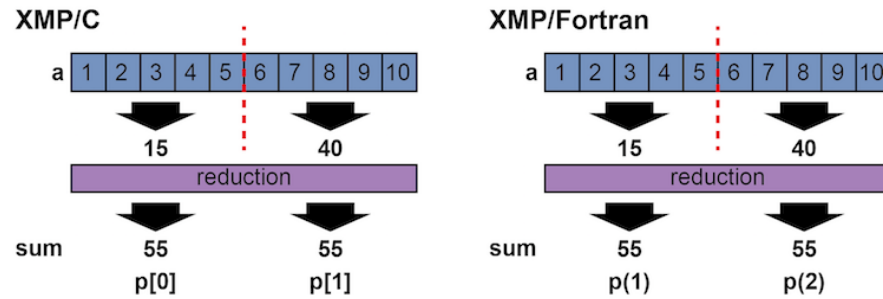
```

    enddo

    write(*,*) sum
15  end program main

```

An operator and target variables for reduction are specified in a `reduction` clause. In the above examples, a “+” operator is specified for the reduction computation to produce a total sum among nodes.



Operations that can be used in a reduction computation are limited to the following associative ones.

C

```

+
*
-
&
5 |
^
&&
||
max
10 min
firstmax
firstmin
lastmax
lastmin

```

Fortran

```

+
*
-
.and.
5 .or.

```

```

    .eqv.
    .neqv.
    max
    min
10  iand
    ior
    ieor
    firstmax
    firstmin
15  lastmax
    lastmin

```

Note: If the reduction variable is a type of floating point, the difference of the order of the executions can make a little bit difference between serial and parallel executions.

3.2.3 Parallelizing Nested Loop

Parallelization of nested loops can be specified similarly for a single one.

```

_____ XcalableMP C _____
#pragma xmp nodes p[2][2]
#pragma xmp template t[10][10]
#pragma xmp distribute t[block][block] onto p

5  int main(){
    int a[10][10];
    #pragma xmp align a[i][j] with t[i][j]

    #pragma xmp loop on t[i][j]
10  for(int i=0;i<10;i++)
    for(int j=0;j<10;j++)
        a[i][j] = i*10+j;

    return 0;
15 }

```

```

_____ XcalableMP Fortran _____
program main
!$xmp nodes p(2,2)
!$xmp template t(10,10)
!$xmp distribute t(block,block) onto p
5  integer :: a(10,10)

```

```

!$xmp align a(j,i) with t(j,i)

!$xmp loop on t(j,i)
  do i=1, 10
    do j=1, 10
10      a(j,i) = i*10+j
    enddo
  enddo
15 end program main

```

3.3 array Construct

The array construct is for work mapping of array assignment statements.

```

_____ XcalableMP C _____
#pragma xmp align a[i] with t[i]
:
#pragma xmp array on t[0:N]
a[0:N] = 1.0;

```

```

_____ XcalableMP Fortran _____
!$xmp align a(i) with t(i)
:
!$xmp array on t(1:N)
a(1:N) = 1.0

```

The above is equivalent to the below.

```

_____ XcalableMP C _____
#pragma xmp align a[i] with t[i]
:
#pragma xmp loop on t[i]
for(int i=0;i<N;i++)
5   a[i] = 1.0;

```

```

_____ XcalableMP Fortran _____
!$xmp align a(i) with t(i)
:
!$xmp loop on t(i)
do i=1, N
5   a(i) = 1.0
enddo

```

This construct can also be applied to multi-dimensional arrays. The triplet notation enables specifying operations for all elements of the array.

```

XcalableMP C
#pragma xmp align a[i][j] with t[i][j]
:
#pragma xmp array on t[:, :]
a[:, :] = 1.0;

```

```

XcalableMP Fortran
!$xmp align a(j,i) with t(j,i)
:
!$xmp array on t(:, :)
a(:, :) = 1.0

```

Note: The template appearing in the `on` clause must have the same shape of arrays in the following statement. The right-hand side value must be identical among all nodes because the array construct is a global (i.e. collective) operation.

4 Data Communication

4.1 shadow Directive and reflect Construct

Stencil computation frequently appears in scientific computations, where array elements $a[i-1]$ and $a[i+1]$ are referenced to update $a[i]$. If $a[i]$ is on the boundary region of a block-distributed array on a node, $a[i+1]$ may reside on another (neighboring) node.

Since it involves large overhead to copy $a[i+1]$ from the neighboring node to update each $a[i]$, a technique of copying collectively elements on the neighboring node to the area added to the distributed array on each node is usually adopted. In XMP, such additional area is called “shadow.”

4.1.1 Declaring Shadow

Shadow areas can be declared with the `shadow` directive. In the example below, an array `a` has shadow areas of width one on both the lower and upper bounds.

```

XcalableMP C
#pragma xmp nodes p[4]
#pragma xmp template t[16]
#pragma xmp distribute t[block] onto p
double a[16];

```

```

5 #pragma xmp align a[i] with t[i]
  #pragma xmp shadow a[1]

```


XcalableMP Fortran

```


!$xmp nodes p(4)
!$xmp template t(16)
!$xmp distribute t(block) onto p
real :: a(16)
5 !$xmp align a(i) with t(i)
!$xmp shadow a(1)

```

XMP/C



XMP/Fortran



In the figure above, colored elements are those that each node owns and white ones are shadow.

Note: Arrays distributed in a cyclic manner cannot have shadow.

In some programs, it is natural that the widths of the shadow area on the lower and upper bounds are different. There is also a case where the shadow area exists only on either of the bounds. In the example below, it is declared that a distributed array *a* has a shadow area of width one only on the upper bound.

XcalableMP C

```

#pragma xmp nodes p[4]
#pragma xmp template t[16]
#pragma xmp distribute t(block) onto p
double a[16];
5 #pragma xmp align a[i] with t[i]
#pragma xmp shadow a[0:1]


```

XcalableMP Fortran

```

!$xmp nodes p(4)
!$xmp template t(16)
!$xmp distribute t(block) onto p
real :: a(16)
5 !$xmp align a(i) with t(i)
!$xmp shadow a(0:1)

```

The values on the left- and right-hand sides of a colon designate the widths on the lower and upper bounds, respectively.

4.1.2 Updating Shadow

To copy data to shadow areas from neighboring nodes, use the `reflect` construct. In the example below, the shadow areas of an array `a` that are of width one on both the upper and lower bounds are updated.

XcalableMP C


```
#pragma xmp reflect (a)

#pragma xmp loop on t[i]
for(int i=1;i<15;i++)
5   a[i] = (a[i-1] + a[i] + a[i+1])/3;
```

XcalableMP Fortran

```
!$xmp reflect (a)

!xmp loop on t(i)
do i=2, 15
5   a(i) = (a(i-1) + a(i) + a(i+1))/3
enddo
```



With this `reflect` directive, in XMP/C, node `p[1]` sends an element `a[4]` to the shadow area on the upper bound on node `p[0]` and `a[7]` to the shadow area on the

lower bound on $p[2]$; $p[0]$ sends an element $a[3]$ to the shadow area on the lower bound on $p[1]$, and $p[2]$ sends $a[8]$ to the shadow area on the upper bound on $p[1]$.


Similarly, in XMP/Fortran, node $p(2)$ sends an element $a(5)$ to the shadow area on the upper bound on node $p(1)$ and $a(8)$ to the shadow area on the lower bound on $p(3)$; $p(1)$ sends an element $a(4)$ to the shadow area on the lower bound on $p(2)$, and $p(3)$ sends $a(9)$ to the shadow area on the upper bound on $p(2)$.

The default behavior of a reflect directive is to update the whole of the shadow area declared by the shadow directive. However, there are some cases where a specific part of the shadow area is to be updated to reduce the communication size at a point of the code.

To update only a specific part of the shadow area, add the `width` clause to the `reflect` directive.

The values on the left- and right-hand sides of a colon in the `width` clause designate the widths on the lower and upper bounds to be updated, respectively. In the example below, only the shadow area on the upper bound is updated.

XcalableMP C
<pre>#pragma xmp reflect (a) width(0:1)</pre>
XcalableMP Fortran
<pre>!\$xmp reflect (a) width(0:1)</pre>



Note: If the widths of the shadow areas to be updated on the upper and lower bounds are equal, that is, for example, `width(1:1)`, you can abbreviate it as `width(1)`.

Note: It is not possible to update the shadow area on a particular node because `reflect` is a collective operation.

The `reflect` directive does not update either the shadow area on the lower bound on the leading node or that on the upper bound on the last node. However, the values

in such areas are needed for stencil computation if the computation needs a periodic boundary condition.

To update such areas, add a `periodic` qualifier into a width clause. Let's look at the following example where an array `a` having shadow areas of width one on both the lower and upper bounds appears.

```


XcalableMP C
#pragma xmp reflect (a) width(/periodic/1:1)

```

```

XcalableMP Fortran
!$xmp reflect (a) width(/periodic/1:1)

```



The `periodic` qualifier has the following effects, in addition to that of a normal `reflect` directive: in XMP/C, node `p[0]` sends an element `a[0]` to the shadow area on the upper bound on node `p[3]`, and `p[3]` sends `a[15]` to the shadow area on the lower bound on `p[0]`; in XMP/Fortran, node `p(1)` sends an element `a(1)` to the shadow area on the upper bound on node `p(4)`, and `p(4)` sends `a(16)` to the shadow area on the lower bound on `p(1)`.

The shadow directive and `reflect` construct can be applied to array distributed in multiple dimensions. The following programs are the examples for two-dimensional distribution.

```

XcalableMP C
#pragma xmp nodes p[3][3]
#pragma xmp template t[9][9]
#pragma xmp distribute t[block][block] onto p
double a[9][9];
5 #pragma xmp align a[i][j] with t[i][j]
#pragma xmp shadow a[1][1]
:
#pragma xmp reflect (a)


```

```

XcalableMP Fortran
!$xmp nodes p(3,3)
!$xmp template t(9,9)
!$xmp distribute t(block,block) onto p
real :: a(9,9)
5 !$xmp align a(j,i) with t(j,i)

```

```
!$xmp shadow a(1,1)
:
!$xmp reflect (a)
```



The central node receives data from the surrounding eight nodes to update its shadow areas. The shadow areas of the other nodes are also updated, which is omitted in the figure.

For some applications, data from ordinal directions are not necessary. In such a case, the data communication from/to the ordinal directions can be avoided by adding an orthogonal clause to a reflect construct.


```
_____ XcalableMP C _____
#pragma xmp reflect (a) orthogonal
```

```
_____ XcalableMP Fortran _____
!$xmp reflect (a) orthogonal
```

Note: The orthogonal clause is effective only for arrays more than one dimension of which is distributed.

Besides, you can also add shadow areas to only specified dimension.

```
_____ XcalableMP C _____
#pragma xmp nodes p[3]
#pragma xmp template t[9]
```



```

#pragma xmp distribute t[block] onto p
double a[9][9];
5 #pragma xmp align a[i][*] with t[i]
  #pragma xmp shadow a[1][0]
  :
  #pragma xmp reflect (a)

```

XcalableMP Fortran

```

!$xmp nodes p[3]
!$xmp template t[9]
!$xmp distribute t[block] onto p
real :: a(9,9)
5 !$xmp align a(*,i) with t(i)
!$xmp shadow a(0,1)
  :
!$xmp reflect (a)


```

For the array `a`, 0 is specified as the shadow width in the dimensions which are not distributed.

4.2 gmove Construct

You can specify a communication for distributed arrays in the form of assignment statements by using the `gmove` construct.

There are three modes of `gmove`; “collective mode,” “in mode,” and “out mode.”



4.2.1 Collective Mode

The copy operation involved by a *collective gmove* is performed collectively, and results in implicit synchronization among the executing nodes.

XscalableMP C

```
#pragma xmp nodes p[4]
#pragma xmp template t[16]
#pragma xmp distribute t[block] onto p
int a[16], b[16];
5 #pragma xmp align a[i] with t[i]
  #pragma xmp align b[i] with t[i]
  :
#pragma xmp gmove
  a[9:5] = b[0:5];
```


XscalableMP Fortran

```
!$xmp nodes p(4)
!$xmp template t(16)
!$xmp distribute t(block) onto p
integer :: a(16), b(16)
5 !$xmp align a(i) with t(i)
  !$xmp align b(i) with t(i)
  :
!$xmp gmove
  a(10:14) = b(1:5)
```

In XMP/C, $p[0]$ sends $b[0]$ - $b[3]$ to $p[2]$ - $p[3]$, and $p[1]$ sends $b[4]$ to $p[3]$. Similarly, in XMP/Fortran, $p(1)$ sends $b(1)$ - $b(4)$ to $p(3)$ - $p(4)$, and $p(2)$ sends $b(5)$ to $p(4)$.

XscalableMP C

```
#pragma xmp nodes p[4]
#pragma xmp template t1[16]
```



```

#pragma xmp template t2[16]
#pragma xmp distribute t1[cyclic] onto p
5 #pragma xmp distribute t2[block] onto p
int a[16], b[16];
#pragma xmp align a[i] with t1[i]
#pragma xmp align b[i] with t2[i]
:
10 #pragma xmp gmove
    a[9:5] = b[0:5];


```

XcalableMP Fortran

```

!$xmp nodes p(4)
!$xmp template t1(16)
!$xmp template t2(16)
!$xmp distribute t1(cyclic) onto p
5 !$xmp distribute t2(block) onto p
integer :: a(16), b(16)
!$xmp align a(i) with t1(i)
!$xmp align b(i) with t2(i)
:
10 !$xmp gmove
    a(10:14) = b(1:5)

```



While array *a* is distributed in a cyclic manner, array *b* is distributed in a block manner.

In XMP/C, *p*[0] sends *b*[0] and *b*[4] to *p*[2] and *p*[3]. *p*[1] sends *b*[1] to *p*[2]. Each element of *p*[2] and *p*[3] will be copied locally. Similarly, in

XMP/Fortran, p(1) sends b(1) and b(5) to p(3) and p(4). p(2) sends b(2) to p(3). Each element of p(3) and p(4) will be copied locally.

By using this method, the shape of a distributed array can be “changed” during computation.


```

XscalableMP C
#pragma xmp nodes p[4]
#pragma xmp template t1[16]
#pragma xmp template t2[16]
int W[4] = {2,4,8,2};
5 #pragma xmp distribute t1[gblock(W)] onto p
  #pragma xmp distribute t2[block] onto p
  int a[16], b[16];
  #pragma xmp align a[i] with t1[i]
  #pragma xmp align b[i] with t2[i]
10 :
  #pragma xmp gmove
    a[:] = b[:];


XscalableMP Fortran
!$xmp nodes p(4)
!$xmp template t1(16)
!$xmp template t2(16)
integer :: W(4) = (/2,4,7,3/)
5 !$xmp distribute t1(gblock(W)) onto p
  !$xmp distribute t2(block) onto p
  integer :: a(16), b(16)
  !$xmp align a(i) with t1(i)
  !$xmp align b(i) with t2(i)
10 :
  !$xmp gmove
    a(:) = b(:)

```

XMP/C



XMP/Fortran



In this example, copying all elements of array **b** that is distributed in a block manner to array **a** that is distributed in a generalized-block manner. For arrays **a** and **b**, communication occurs if corresponding elements reside in different nodes (arrows illustrate communication between nodes in the figures).

In an assignment statement, if a scalar (i.e. one element of an array or a variable) is specified on the right-hand side and an array section are specified on the left-hand side, the operation will be broadcast communication.


XcalableMP C

```
#pragma xmp nodes p[4]
#pragma xmp template t[16]
#pragma xmp distribute t[block] onto p
int a[16], b[16];
5 #pragma xmp align a[i] with t[i]
  #pragma xmp align b[i] with t[i]
  :
#pragma xmp gmove
  a[9:5] = b[0];
```


XcalableMP Fortran

```
!$xmp nodes p(4)
!$xmp template t(16)
!$xmp distribute t(block) onto p
integer :: a(16), b(16)
5 !$xmp align a(i) with t(i)
  !$xmp align b(i) with t(i)
  :
!$xmp gmove
  a(10:14) = b(1)
```

XMP/C



XMP/Fortran



In this example, in XMP/C, an array element $b[0]$ of node $p[0]$ will be broadcasted to the specified array section on node $p[2]$ and $p[3]$. Similarly, in XMP/Fortran, an array element $b(1)$ of node $p(1)$ will be broadcasted to the specified array section on node $p(3)$ and $p(4)$.

Not only distributed arrays but also replicated arrays can be specified on the right-hand side.

XcalableMP C

```
#pragma xmp nodes p[4]
#pragma xmp template t[16]
#pragma xmp distribute t[block] onto p
int a[16], b[16], c;
```

```

5  #pragma xmp align a[i] with t[i]
    :
    #pragma xmp gmove
    a[9:5] = b[0:5];

```

XscalableMP Fortran

```

!$xmp nodes p(4)
!$xmp template t(16)
!$xmp distribute t(block) onto p
integer :: a(16), b(16), c
5  !$xmp align a(i) with t(i)
    :
!$xmp gmove
    a(10:14) = b(1:5)

```

In this example, a replicated array `b` is locally copied to distributed array `a` without communication.

XscalableMP C

```

#pragma xmp nodes p[4]
#pragma xmp template t1[8]
#pragma xmp template t2[16]
#pragma xmp distribute t1[block] onto p
5  #pragma xmp distribute t2[block] onto p
int a[8][16], b[8][16];
#pragma xmp align a[i][*] with t1[i]
#pragma xmp align b[*][i] with t2[i]
    :
10 #pragma xmp gmove
    a[0][:] = b[0][:];

```



XscalableMP Fortran

```

!$xmp nodes p(4)
!$xmp template t1(8)
!$xmp template t2(16)
!$xmp distribute t1(block) onto p
5  !$xmp distribute t2(block) onto p
integer :: a(16,8), b(8,16)
!$xmp align a(*,i) with t1(i)
!$xmp align b(i,*) with t2(i)
    :
10 #pragma xmp gmove
    a(:,1) = b(:,1)

```

In this example, in XMP/C, `b[0][0:2]` on `p[0]`, `b[0][2:2]` of `p[1]`, `b[0][4:2]` on `p[2]` and `b[0][6:2]` on `p[3]` are copied to `a[0][:]` on `p[0]`. Similarly, in XMP/Fortran, `b(1:2,1)` on `p(1)`, `b(3:4,1)` of `p(2)`, `b(5:6,1)` on `p(3)` and `b(7:8,1)` on `p(4)` are copied to `a(:,1)` on `p(1)`.

XMP/C**XMP/Fortran****4.2.2 In Mode**

The right-hand side data of the assignment, all or part of which may reside outside the executing node set, can be transferred from its owner nodes to the executing nodes by an *in gmove*.

```

XcalableMP C
#pragma xmp nodes p[4]
#pragma xmp template t[4]
#pragma xmp distribute t[block] onto p
double a[4], b[4];
5 #pragma xmp align a[i] with t[i]
  #pragma xmp align b[i] with t[i]
  :
  #pragma xmp task on p[0:2]
  #pragma xmp gmove in
10   a[0:2] = b[2:2]
  #pragma xmp end task

XcalableMP Fortran
!$xmp nodes p(4)
!$xmp template t(4)
!$xmp distribute t(block) onto p
real :: a(4), b(4)
5 !$xmp align a(i) with t(i)
!$xmp align b(i) with t(i)

```


```

:
!$xmp task on p(1:2)
!$xmp gmove in
10   a(1:2) = b(3:4)
!$xmp end task


```

In this example, the `task` directive divides four nodes into two sets, the first-half and the second-half. A `gmove` construct that is in an *in* mode copies data using a *get* operation from the second-half node to the first-half node.

XMP/C



XMP/Fortran



4.2.3 Out Mode

The left-hand side data of the assignment, all or part of which may reside outside the executing node set, can be transferred from the executing nodes to its owner nodes by an *out* `gmove` construct.

```

XscalableMP C
#pragma xmp nodes p[4]
#pragma xmp template t[4]

```

```

#pragma xmp distribute t[block] onto p
double a[4], b[4];
5 #pragma xmp align a[i] with t[i]
  #pragma xmp align b[i] with t[i]
  :
  #pragma xmp task on p[0:2]
  #pragma xmp gmove out
10   b[2:2] = a[0:2]
  #pragma xmp end task

```

XcalableMP Fortran

```

!$xmp nodes p(4)
!$xmp template t(4)
!$xmp distribute t(block) onto p
real :: a(4), b(4)
5 !$xmp align a(i) with t(i)
  !$xmp align b(i) with t(i)
  :
  !$xmp task on p(1:2)
  !$xmp gmove out
10   b(3:4) = a(1:2)
  !$xmp end task

```

A `gmove` construct that is in *out* mode copies data using a *put* communication from the first-half nodes to the second-half nodes.

4.3 barrier Construct

The barrier construct executes a barrier synchronization.

XcalableMP C

```
#pragma xmp barrier
```

XcalableMP Fortran

```
!$xmp barrier
```



You can specify a node set on which the barrier synchronization is to be performed by using the `on` clause. In the below example, a barrier synchronization is performed among the first two nodes of `p`.

XcalableMP C

```
#pragma xmp barrier on p[0:2]
```

XcalableMP Fortran

```
!$xmp barrier on p(1:2)
```

XMP/C**XMP/Fortran****4.4 reduction Construct**

This construct performs a *reduction* operation. It has the same meaning as the *reduction* clause of the *loop* construct, but this construct can be specified anywhere as an *executable* construct.

XcalableMP C

```
#pragma xmp nodes p[4]
:
sum = xmpc_node_num() + 1;
#pragma xmp reduction (+:sum)
```

XcalableMP Fortran

```
!$xmp nodes p(4)
:
sum = xmp_node_num()
!$xmp reduction (+:sum)
```

You can specify the executing node set by using the *on* clause. In the below example, only the values on the last two of the four nodes are targeted by the *reduction* construct.

XMP/C

	p[0]	p[1]	p[2]	p[3]	
sum =	1	2	3	4	reduction(+:sum)
sum =	10	10	10	10	

XMP/Fortran

	p(1)	p(2)	p(3)	p(4)	
sum =	1	2	3	4	reduction(+:sum)
sum =	10	10	10	10	

XcalableMP C

```
#pragma xmp nodes p[4]
:
sum = xmpc_node_num() + 1;
#pragma xmp reduction (+:sum) on p[2:2]
```

XcalableMP Fortran

```
!$xmp nodes p(4)
:
sum = xmp_node_num()
!$xmp reduction (+:sum) on p(3:4)
```

The operators you can use in the reduction construct are as follows:

XcalableMP C

```
+
*
-
&
5 |
^
&&
||
max
10 min
```

XcalableMP Fortran

```
+
*
-
```

XMP/C

	p[0]	p[1]	p[2]	p[3]	
sum =	1	2	3	4	
	<hr/>				
sum =	1	2	7	7	reduction (+:sum) on p[2:2]

XMP/Fortran

	p(1)	p(2)	p(3)	p(4)	
sum =	1	2	3	4	
	<hr/>				
sum =	1	2	7	7	reduction (+:sum) on p(3:4)

```

5 .and.
  .or.
  .eqv.
  .neqv.
max
min
10 iand
   ior
   ieor

```

Note: In contrast to the `reduction` clause of the `loop` construct, which precedes loops, the `reduction` construct does not accept operators of `firstmax`, `firstmin`, `lastmax`, and `lastmin`.

Note: Similar to the `reduction` clause, the `reduction` construct may generate slightly different results in a parallel execution from in a sequential execution, because the results depends on the order of associating the value.

4.5 bcast Construct

The `bcast` construct broadcasts the values of the variables on the node specified by the `from` clause, that is, the *root node*, to the node set specified by the `on` clause. If there is no `from` clause, the first node of the executing node set is selected as the root node. If there is no `on` clause, the current executing node set of the construct is selected as the executing node set.

In the below example, the first node of the node set `p` is the root node.

```

XcalableMP C
#pragma xmp nodes p[4]
:
num = xmpc_node_num() + 1;
#pragma xmp bcast (num)

```

```

XcalableMP Fortran
!$xmp nodes p(4)
:
num = xmp_node_num()
!$xmp bcast (num)

```

XMP/C

	p[0]	p[1]	p[2]	p[3]	
num =	1	2	3	4	bcast (num)
num =	1	1	1	1	

XMP/Fortran

	p(1)	p(2)	p(3)	p(4)	
num =	1	2	3	4	bcast (num)
num =	1	1	1	1	

In the below example, the last node is the `from` clause.

```

XcalableMP C
#pragma xmp nodes p[4]
:
num = xmpc_node_num() + 1;
#pragma xmp bcast (num) from p[3]

```

XcalableMP Fortran

```
!$xmp nodes p(4)
:
num = xmp_node_num()
!$xmp bcast (num) from p(4)
```

XMP/C

	p[0]	p[1]	p[2]	p[3]	
num =	1	2	3	4	
					bcast (num) from p[3]
num =	4	4	4	4	

XMP/Fortran

	p(1)	p(2)	p(3)	p(4)	
num =	1	2	3	4	
					bcast (num) from p(4)
num =	4	4	4	4	

In the below example, only the last three of four nodes are the executing node set of the bcast construct.

XcalableMP C

```
#pragma xmp nodes p[4]
:
sum = xmpc_node_num() + 1;
#pragma xmp bcast (num) from p[3] on p[1:3]
```

XcalableMP Fortran

```
!$xmp nodes p(4)
:
sum = xmp_node_num()
!$xmp bcast (num) from p(4) on p(2:4)
```

4.6 wait_async Construct

Communication directives (i.e. `reflect`, `gmove`, `reduction`, `bcast`, `reduce_shadow`) can perform asynchronous communication if the `async` clause is added. The

XMP/C

	p[0]	p[1]	p[2]	p[3]	
num =	1	2	3	4	
	<hr/>				bcast (num) from p[3] on p[1:3]
num =	1	4	4	4	

XMP/Fortran

	p(1)	p(2)	p(3)	p(4)	
num =	1	2	3	4	
	<hr/>				bcast (num) from p(4) on p(2:4)
num =	1	4	4	4	

`wait_async` construct is used to guarantee the completion of asynchronous communication.

	XcalableMP C
	<code>#pragma xmp bcast (num) async(1)</code>
	<code>:</code>
	<code>#pragma xmp wait_async (1)</code>

	XcalableMP Fortran
	<code>!\$xmp bcast (num) async(1)</code>
	<code>:</code>
	<code>!\$xmp wait_async (1)</code>

Since the `bcast` directive has an `async` clause, communication may not be completed immediately after the `bcast` directive. The completion of that communication is guaranteed with the `wait_async` construct having the same value as that of the `async` clause. Therefore, between the `bcast` construct and the `wait_async` constructs, you may not reference the target variable of the `bcast` directive.

Hint: By performing computation without a dependency relationship with the variable specified by the `bcast` directive after the `bcast` directive, overlap of communication and computation can be performed, so the total computation time may be small.

Note: Expressions that can be specified in the `async` clause are of type `int`, in XMP/C, or `integer`, in XMP/Fortran.

4.7 reduce_shadow Construct

The `reduce_shadow` directive adds the value of a shadow object to the corresponding data object of the array.

```

XcalableMP C
#pragma xmp nodes p[2]
#pragma xmp template t[8]
#pragma xmp distribute t[block] onto p
int a[8];
5 #pragma xmp align a[i] with t[i]
  #pragma xmp shadow a[1]
  :
  #pragma xmp loop on t[i]
    for(int i=0;i<8;i++)
10     a[i] = i+1;

#pragma xmp reflect (a)
#pragma xmp reduce_shadow (a)

```

```

XcalableMP Fortran
!$xmp nodes p(2)
!$xmp template t(8)
!$xmp distribute t(block) onto p
integer a(8)
5 !$xmp align a(i) with t(i)
!$xmp shadow a(1)



!$xmp loop on t(i)
do i=1, 8
10   a(i) = i
enddo

!$xmp reflect (a)
!$xmp reduce_shadow (a)

```

For the above example, in XMP/C, `a[3]` on `p[0]` has a value of eight, and `a[4]` on `p[1]` has a value of ten. Similarly, in XMP/Fortran, `a(4)` of `p(1)` has a value of eight, and `a(5)` on `p(2)` has a value of ten.

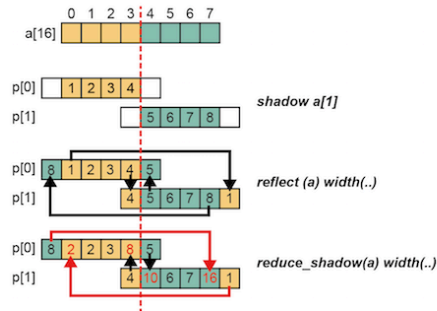
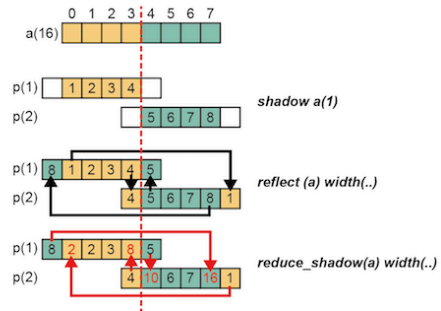
You can add the periodic modifier to the width clause to update the shadow area periodically.

XMP/C**XMP/fortran****XcalableMP C**

```
#pragma xmp reflect (a) width(/periodic/1)
#pragma xmp reduce_shadow (a) width(/periodic/1)
```

XcalableMP Fortran

```
!$xmp reflect (a) width(/periodic/1)
!$xmp reduce_shadow (a) width(/periodic/1)
```

XMP/C**XMP/fortran**

In addition to the first example, in XMP/C, `a[0]` on `p[0]` has a value of two, and `a[7]` on `p[1]` has a value of 16. Similarly, in XMP/fortran, `a(1)` in `p(1)` has a value of two, and `a(8)` in `p(2)` has a value of 16.

5 Local-view Programming

5.1 Introduction

The user uses Coarray in the local-view model to describe one-sided communication. In XMP, put/get communication and some synchronization functions are supported.

If the target system supports Remote Direct Memory Access (RDMA) in the hardware, one-sided communication in the local-view model can achieve better performance compared to the global-view model. However, it requires more effort to describe parallel program since all communication should be specified in detail.

Coarray in XMP/Fortran is upward-compatible with that in Fortran 2008; coarray in XMP/C is defined as an extension to the base language.

An execution entity in a local-view XMP program is referred to as an “image” while a “node” in a global-view. These two words have almost the same meaning in XMP.

5.2 Coarray Declaration

_____ XcalableMP C _____
<code>int a[10]:[*];</code>
_____ XcalableMP Fortran _____
<code>integer a(10)[*]</code>

In XMP/C, the user declares a coarray by adding “:[*]” after the array declaration. In XMP/Fortran, the user declares a coarray by adding “[*]” after the array declaration.

Note: Based on Fortran 2008, coarrays should have the same size among all images.

A coarray can be accessed by remote images with assignment statements. Of course, it can be also accessed by the local image like ordinary arrays.

5.3 Put Communication

When a coarray appears in the left-hand side of an assignment statement, it involves *put* communication.

_____ XcalableMP C _____
<code>int a[10]:[*], b[10];</code>

```

if (xmpc_this_image() == 0)
  a[0:3]:[1] = b[3:3];

```

XcalableMP Fortran

```

integer a(10)[*]
integer b(10)


if (this_image() == 1) then
  a(1:3)[2] = b(3:5)
end if

```


The integer in the square bracket specifies the target image index. The image index is 0-based, in XMP/C, or 1-based, in XMP/Fortran. `xmpc_this_image()` in XMP/C and `this_image()` in XMP/Fortran return the current image index.

In the above example, in XMP/C, an image 0 puts `b[3:3]` to `a[0:3]` on image 1; in XMP/Fortran, an image 1 puts `b(3:5)` to `a(1:3)` on image 2. The following figure illustrates the one-sided communication performed in the example.

XMP/C



XMP/Fortran



5.4 Get Communication

When a coarray appears in the right-hand side of an assignment statement, it involves *get* communication.

XcalableMP C

```

int a[10]:[*], b[10];

if (xmpc_this_image() == 0)
  b[3:3] = a[0:3]:[1];

```

XcalableMP Fortran

```

integer a(10)[*]
integer b(10)


```

```

if (this_image() == 1) then
  b(3:5) = a(1:3)[2]
end if

```

In the above example, in XMP/C, an image 0 gets $a[0:3]$ from an image 1 and copies it to $b[3:3]$; in XMP/Fortran, an image 1 gets $a(1:3)$ from an image 2 and copies it to $b(3:5)$ of an image 1. The following figure illustrates the one-sided communication performed in the example.



Hint: As illustrated above, get communication involves an extra step to send a request to the target node. Put communication achieves better performance than get because there is no such extra step.

5.5 Synchronization

Here, we introduce “sync all,” which is most frequently used among synchronization features for coarrays.

XcalableMP C

```

void xmp_sync_all(int *status)

```


XcalableMP Fortran

```

sync all

```

At “sync all,” each image waits until all issued one-sided communication is complete and then performs barrier synchronization among the entire images.



In the above example, the left image puts data to the right image and both nodes invoke `sync all`. When both nodes return from it, the execution continues to the following statements.

5.5.1 `sync all`

	XcalableMP C	
void	<code>xmp_sync_all(int *status)</code>	
	XcalableMP Fortran	
	<code>sync all</code>	

Each image waits until all one-sided communication issued is complete, and performs barrier synchronization.

5.5.2 `sync images`

	XcalableMP C	
void	<code>xmp_sync_images(int num, int *image-set, int *status)</code>	
	XcalableMP Fortran	
	<code>sync images (image-set)</code>	

Each image in the specified images waits until all one-sided communication issued is complete, and performs barrier synchronization among the images.

```

XcalableMP C
int image_set[3] = {0,1,2};
xmp_sync_images(3, image_set, NULL);

```

```

XcalableMP Fortran
integer :: image_set(3) = (/ 1, 2, 3/)
sync images (image_set)

```

5.5.3 sync memory

```

XcalableMP C
void xmp_sync_memory(int *status)

```

```

XcalableMP Fortran
sync memory

```

Each image waits until all one-sided communication is complete. This function/s-statement does not imply barrier synchronization, unlike `sync all` and `sync images`, and therefore can be locally executed.

5.6 post/wait Construct

5.7 lock/unlock Construct

6 Procedure Interface

Procedure calls in XMP are the same as the base language. Procedure calls between different languages and external library are also possible if the base language supports them.

In the below example, `sub1()` calls `sub2()` with a distributed array as an argument.

```

XcalableMP C
void sub1(){
  #pragma xmp nodes p[2]
  #pragma xmp template t[10]
  #pragma xmp distribute t[block] onto p
5   double x[10];
  #pragma xmp align x[i] with t[i]
    sub2(x);
}
10 void sub2(double a[10]){

```

```

#pragma xmp nodes p[2]
#pragma xmp template t[10]
#pragma xmp distribute t[block] onto p
    double a[10];
15 #pragma xmp align a[i] with t[i]
    :
    }

```

XcalableMP Fortran

```

subroutine sub1()
!$xmp nodes p(2)
!$xmp template t(10)
!$xmp distribute t(block) onto p
5   real x(10)
!$xmp align x(i) with t(i)
    call sub2(x)
end subroutine

10 subroutine sub2(a)
!$xmp nodes p(2)
!$xmp template t(10)
!$xmp distribute t(block) onto p
    real a(10)
15 !$xmp align a(i) with t(i)
    :
end subroutine

```

If you want to use distributed arrays in arguments as distributed arrays in the called procedure, you need to redefine the shape of the distributed array in the procedure.

But, if you want to use the distributed array in the argument as a duplicate array in the called procedure, you do not need to redefine them.



XcalableMP C

```

void sub1(){
#pragma xmp nodes p[2]
#pragma xmp template t[10]
#pragma xmp distribute t[block] onto p
5   double x[10];
#pragma xmp align x[i] with t[i]
    sub2(x);
}

10 void sub2(double a[5]){
    :
}

```

XMP/C**XMP/Fortran**

XscalableMP Fortran

```



subroutine sub1()
!$xmp nodes p(2)
!$xmp template t(10)
!$xmp distribute t(block) onto p
5   real x(10)
!$xmp align x(i) with t(i)
    call sub2(x)
end subroutine

10 subroutine sub2(a)
    real a(5)
    :
end subroutine

```

References

1. Robert W. Numrich and John Reid, "Co-Array Fortran for parallel programming", ACM SIGPLAN Fortran Forum, Vol. 17, No. 2 (1998).
2. UPC Consortium, "UPC Specifications, v1.2", Lawrence Berkeley National Lab (LBNL-59208) (2005).

XMP/C**XMP/Fortran**

- David Callahan, Bradford L. Chamberlain and Hans P. Zima, "The Cascade High Productivity Language", Proc. 9th Int'l. Workshop on High-Level Parallel Programming Models and Supportive Environments (HIPS 2004), pp. 52–60 (2004).